

Lecture 11 — Planning

TDT4136: Introduction to Artificial Intelligence

Xavier F. C. Sánchez Díaz

Department of Computer Science
Faculty of Information Technology and Electrical Engineering
Norwegian University of Science and Technology

October 31, 2024



Outline

- 1 What is Planning?
- 2 Representing the world
- 3 PDDL
- 4 How to plan?
 - Forward planning
 - Backward search
 - Partial ordering
- 5 Heuristic Planning
- 6 Planning in Complex Environments

What is Planning?

Classical Planning

Find a sequence of actions to accomplish a goal in an environment that is:

- ▶ discrete
- ▶ deterministic
- ▶ static
- ▶ fully observable

What is Planning?

Classical Planning

Find a sequence of actions to accomplish a goal in an environment that is:

- ▶ discrete
- ▶ deterministic
- ▶ static
- ▶ fully observable

In other words, what we have discussed before while studying **search**!

Why do we plan?

What is Planning?

- ▶ To address a new situation or problem

Why do we plan?

What is Planning?

- ▶ To address a new situation or problem
- ▶ When the task is complex

Why do we plan?

What is Planning?

- ▶ To address a new situation or problem
- ▶ When the task is complex
- ▶ when the environment imposes a **high risk or cost!**

Why do we plan?

What is Planning?

- ▶ To address a new situation or problem
- ▶ When the task is complex
- ▶ when the environment imposes a **high risk or cost!**
- ▶ when collaborating with others

Why do we plan?

What is Planning?

- ▶ To address a new situation or problem
- ▶ When the task is complex
- ▶ when the environment imposes a **high risk or cost!**
- ▶ when collaborating with others

Why do we plan?

What is Planning?

- ▶ To address a new situation or problem
- ▶ When the task is complex
- ▶ when the environment imposes a **high risk or cost!**
- ▶ when collaborating with others

We *explicitly* plan when it is strictly necessary.

How do we do it?

What is Planning?

There are multiple ways to define a planning problem, and several ways of finding a valid plan.

How do we do it?

What is Planning?

There are multiple ways to define a planning problem, and several ways of finding a valid plan.

So far, we have covered:

- ▶ Problem-solving by **searching**

How do we do it?

What is Planning?


There are multiple ways to define a planning problem, and several ways of finding a valid plan.

So far, we have covered:

- ▶ Problem-solving by **searching**
- ▶ **Logic** and **satisfiability**

Representing the world

Representing the world



The **state of the world** can be described in different ways.

An example: Wumpus World

Representing the world

- ▶ A **partially observable** world, with **sensors** and a limited set of actions
- ▶ We act rationally by updating **our belief** of the world
- ▶ The world is **stored as facts in a knowledge base** (a logical agent can solve this!)

<https://thiagodnf.github.io/wumpus-world-simulator/>

What happens when we move?

What happens when we move?

Time is important!

Dealing with time(steps)

Representing the world

- Some aspects of the world change from time to time. We call them **fluents**.

Dealing with time(steps)

Representing the world

- ▶ Some aspects of the world change from time to time. We call them **fluents**.
- ▶ For example: the state of a grid cell—‘free’ or ‘occupied’ at time 0

Dealing with time(steps)

Representing the world

- ▶ Some aspects of the world change from time to time. We call them **fluents**.
- ▶ For example: the state of a grid cell—‘free’ or ‘occupied’ at time 0
- ▶ (First Order) Logic cannot deal with it

Dealing with time(steps)

Representing the world

- ▶ Some aspects of the world change from time to time. We call them **fluents**.
- ▶ For example: the state of a grid cell—‘free’ or ‘occupied’ at time 0
- ▶ (First Order) Logic cannot deal with it
 - ▶ There are extensions that can handle it, like **linear temporal** or **computational tree** logics!

Dealing with time(steps)

Representing the world

- ▶ Some aspects of the world change from time to time. We call them **fluents**.
- ▶ For example: the state of a grid cell—‘free’ or ‘occupied’ at time 0
- ▶ (First Order) Logic cannot deal with it
 - ▶ There are extensions that can handle it, like **linear temporal** or **computational tree** logics!
- ▶ An agent's actions can change aspects of the world (fluents) but not all

Dealing with time(steps)

Representing the world

- ▶ Some aspects of the world change from time to time. We call them **fluents**.
- ▶ For example: the state of a grid cell—‘free’ or ‘occupied’ at time 0
- ▶ (First Order) Logic cannot deal with it
 - ▶ There are extensions that can handle it, like **linear temporal** or **computational tree** logics!
- ▶ An agent’s actions can change aspects of the world (fluents) but not all
- ▶ The agent needs to keep track of fluents, and know what remains unchanged!

A snapshot of the world

Representing the world

Time 0

I am at *cellA1* facing *east*, I feel a *breeze* and have *1 arrow*.

A snapshot of the world

Representing the world

Time 0

I am at *cellA1* facing *east*, I feel a *breeze* and have 1 *arrow*.

If we decide to move *Forward*, then in logic:

$$Location_{cellA1}^0 \wedge FacingEast^0 \wedge Forward^0 \implies Location_{cellA2}^1 \wedge \neg Location_{cellA1}^1$$

... and although the *arrows* and *breeze* percepts were not modelled, we would have 4 *directions* $\times T$ time steps $\times n^2$ locations.

A snapshot of the world

Representing the world

Time 0

I am at *cellA1* facing *east*, I feel a *breeze* and have 1 *arrow*.

If we decide to move *Forward*, then in logic:

$$Location_{cellA1}^0 \wedge FacingEast^0 \wedge Forward^0 \implies Location_{cellA2}^1 \wedge \neg Location_{cellA1}^1$$

... and although the *arrows* and *breeze* percepts were not modelled, we would have 4 *directions* $\times T$ time steps $\times n^2$ locations.

It is extremely expensive and inefficient!

Planning Domain Definition Language

PDDL

Instead, we can use **PDDL**, which uses *actions* with *preconditions* and **effects**:

Planning Domain Definition Language

PDDL

Instead, we can use **PDDL**, which uses **actions** with **preconditions** and **effects**:

Action(Move(who, from, to))

Precond : $At(who, from) \wedge Adj(from, to) \wedge \neg Pit(to)$

Effect : $\neg At(who, from) \wedge At(who, to)$

Planning Domain Definition Language

PDDL

Action(Move(who, from, to))

Precond : At(who, from) \wedge Adj(from, to) \wedge \neg Pit(to)

Effect : \neg At(who, from) \wedge At(who, to)

- ▶ *Move* is the **action** being defined
- ▶ *who*, *from* and *to* are **variables**
- ▶ *Precond* describes the **state** of the world needed for the action *to occur*
- ▶ *Effect* describes the **resulting state** after acting

States in PDDL

PDDL

- The world is closed—any **fluents** not mentioned are *False*

States in PDDL

PDDL

- ▶ The world is closed—any *fluents* not mentioned are *False*
- ▶ Unique names—different literals refer to different entities (like *cellA1* and *cellA2*)

States in PDDL

PDDL

- ▶ The world is closed—any *fluents* not mentioned are *False*
- ▶ Unique names—different literals refer to different entities (like *cellA1* and *cellA2*)
- ▶ No uncertain or negated literals

States in PDDL

PDDL

- ▶ The world is closed—any *fluents* not mentioned are *False*
- ▶ Unique names—different literals refer to different entities (like *cellA1* and *cellA2*)
- ▶ No uncertain or negated literals

States in PDDL


PDDL

- ▶ The world is closed—any **fluents** not mentioned are *False*
- ▶ Unique names—different literals refer to different entities (like *cellA1* and *cellA2*)
- ▶ No uncertain or negated literals


As with search, we also need **starting** and **goal** states.

Example: Block world

PDDL




Start State



Goal State

What would the **start and goal** states look like in **PDDL**?


Example: Block world



- **Start:** *On(A, Table), On(B, Table), On(C, A), Clear(B), Clear(C)*¹
- **Goal:** *On(A, B), On(B, C)*

¹Commas are a shorthand for *ANDs* (\wedge) and semicolons for *ORs* (\vee)

Example: Block world




- **Start:** *On(A, Table), On(B, Table), On(C, A), Clear(B), Clear(C)*¹
- **Goal:** *On(A, B), On(B, C)*

What would the **actions** look like in **PDDL**?

¹Commas are a shorthand for **ANDs** (\wedge) and semicolons for **ORs** (\vee)

Example: Block world



- **Start:** $On(A, Table), On(B, Table), On(C, A), Clear(B), Clear(C)$
- **Goal:** $On(A, B), On(B, C)$

Action(Move(block, x, y)

PRECOND : On(block, x), Clear(block), Clear(y), Block(block), Block(y)

EFFECT : On(block, y), Clear(x), $\neg On(block, x)$, $\neg Clear(y)$)

Notice how any variable in the **effect** must appear in the **precondition**!

Example: Block world

- **Start:** $On(A, Table), On(B, Table), On(C, A), Clear(B), Clear(C)$
- **Goal:** $On(A, B), On(B, C)$



Action(Move(block, x, y)

PRECOND : On(block, x), Clear(block), Clear(y)


EFFECT : On(block, y), Clear(x), $\neg On(block, x)$, $\neg Clear(y)$)

Action(MoveToTable(block, x)

PRECOND : On(block, x), Clear(block), Clear(Table)

EFFECT : On(block, Table), Clear(x), $\neg On(block, x)$)

Example: Block world



Solution: [*MoveToTable(C, A)*, *Move(B, Table, C)*, *Move(A, Table, B)*]

Adding and deleting

PDDL

Action(Move(block, x, y)

PRECOND : On(block, x), Clear(block), Clear(y)

EFFECT : On(block, y), Clear(x), \neg On(block, x), \neg Clear(y))

At **each step** of the plan, the **state of the world** is **modified** depending on the **effect** of the action taken:

Adding and deleting

PDDL

Action(Move(block, x, y)

PRECOND : On(block, x), Clear(block), Clear(y)

EFFECT : On(block, y), Clear(x), \neg On(block, x), \neg Clear(y))

At **each step** of the plan, the **state of the world** is **modified** depending on the **effect** of the action taken:

- Positive fluents are **added**

Adding and deleting

PDDL

Action(Move(block, x, y)

PRECOND : On(block, x), Clear(block), Clear(y)

EFFECT : On(block, y), Clear(x), \neg On(block, x), \neg Clear(y))

At **each step** of the plan, the **state of the world** is **modified** depending on the **effect** of the action taken:

- ▶ Positive fluents are **added**
- ▶ Negative fluents are **deleted**

Adding and deleting

PDDL

Action(Move(block, x, y)

PRECOND : On(block, x), Clear(block), Clear(y)

EFFECT : On(block, y), Clear(x), \neg On(block, x), \neg Clear(y))

At **each step** of the plan, the **state of the world** is **modified** depending on the **effect** of the action taken:

- ▶ Positive fluents are **added**
- ▶ Negative fluents are **deleted**

Adding and deleting

PDDL

Action(Move(block, x, y)

PRECOND : On(block, x), Clear(block), Clear(y)

EFFECT : On(block, y), Clear(x), \neg On(block, x), \neg Clear(y))

At **each step** of the plan, the **state of the world** is **modified** depending on the **effect** of the action taken:

- ▶ Positive fluents are **added**
- ▶ Negative fluents are **deleted**

These are known as the ADD and DEL lists, and allow us to calculate the state **s** at the next time step after taking action **a**:

$$s^{(t+1)} = (s^{(t)} \setminus DEL(a)) \cup ADD(a)$$

Designing actions

PDDL

There are some other things to consider.

- How would the system differentiate between *MoveToTable(block, x)* and *Move(block, x, Table)*?

Designing actions

PDDL

There are some other things to consider.

- ▶ How would the system differentiate between *MoveToTable(block, x)* and *Move(block, x, Table)*?
 - ▶ Larger-than-necessary search space!

Designing actions

PDDL

There are some other things to consider.

- ▶ How would the system differentiate between *MoveToTable(block, x)* and *Move(block, x, Table)*?
 - ▶ Larger-than-necessary search space!
- ▶ Some operators are a bit dumb, i.e., *Move(B, C, C)* as it introduces an inconsistent *Clear(C), \neg Clear(C)*

Designing actions

PDDL

There are some other things to consider.

- ▶ How would the system differentiate between *MoveToTable(block, x)* and *Move(block, x, Table)*?
 - ▶ Larger-than-necessary search space!
- ▶ Some operators are a bit dumb, i.e., *Move(B, C, C)* as it introduces an inconsistent *Clear(C), \neg Clear(C)*

Designing actions

PDDL

There are some other things to consider.

- ▶ How would the system differentiate between *MoveToTable(block, x)* and *Move(block, x, Table)*?
 - ▶ Larger-than-necessary search space!
- ▶ Some operators are a bit dumb, i.e., *Move(B, C, C)* as it introduces an inconsistent *Clear(C), \neg Clear(C)*
- ▶ A solution is to use the predicate *Block* so that in order to *Move(x, y, z)* we need *Block(x), Block(y)*

Designing actions

PDDL

There are some other things to consider.

- ▶ How would the system differentiate between *MoveToTable(block, x)* and *Move(block, x, Table)*?
 - ▶ Larger-than-necessary search space!
- ▶ Some operators are a bit dumb, i.e., *Move(B, C, C)* as it introduces an inconsistent *Clear(C), $\neg\text{Clear}(C)$*
- ▶ A solution is to use the predicate *Block* so that in order to *Move(x, y, z)* we need *Block(x), Block(y)*
- ▶ Another approach is to use $\neg(y = z)$ as a precondition for *Move(x, y, z)*

Designing actions

PDDL

There are some other things to consider.

- ▶ How would the system differentiate between *MoveToTable(block, x)* and *Move(block, x, Table)*?
 - ▶ Larger-than-necessary search space!
- ▶ Some operators are a bit dumb, i.e., *Move(B, C, C)* as it introduces an inconsistent *Clear(C), \neg Clear(C)*
- ▶ A solution is to use the predicate *Block* so that in order to *Move(x, y, z)* we need *Block(x), Block(y)*
- ▶ Another approach is to use $\neg(y = z)$ as a precondition for *Move(x, y, z)*

Designing actions

PDDL

There are some other things to consider.

- ▶ How would the system differentiate between *MoveToTable(block, x)* and *Move(block, x, Table)*?
 - ▶ Larger-than-necessary search space!
- ▶ Some operators are a bit dumb, i.e., *Move(B, C, C)* as it introduces an inconsistent *Clear(C), \neg Clear(C)*
- ▶ A solution is to use the predicate *Block* so that in order to *Move(x, y, z)* we need *Block(x), Block(y)*
- ▶ Another approach is to use $\neg(y = z)$ as a precondition for *Move(x, y, z)*

It is **not easy!**

How to plan?

There are several ways to come up with a **feasible** plan, and the **search space** might be different on each approach.

How to plan?

There are several ways to come up with a **feasible** plan, and the **search space** might be different on each approach.

- ▶ **State-space planning**: search through nodes representing states of the world. A plan is a **path** through the space
- ▶ **Plan-space planning**: search through partially instantiated operators and constraints—starts with a partial, *possibly incorrect* plan and then apply changes to correct it.
- ▶ **Heuristic planning**: search for a sequence of actions and evaluate your plan using an objective function.


Algorithms for classical planning

How to plan?

- ▶ **Forward** (progression) search
- ▶ **Backward** (regression) search
- ▶ **Logical Inference**

Forward planning


How to plan?



1. Determine all **actions** applicable
2. **Ground**² the actions by replacing any **variable** with **constants**
3. Choose an **action** to apply
4. Determine the new state of the world and update the knowledge based according to the action description
5. Repeat this process until the **goal state** is reached


²Instantiate, if you will

Forward planning




What is your plan?

Forward planning



How many possible **first actions** are there?

Forward planning



How many possible **first actions** are there?

How do we know which one is the *best* one?

Branching Factor

Forward Search

- **Forward search** can have a **very large** branching factor

Branching Factor

Forward Search

- ▶ **Forward search** can have a **very large** branching factor
 - ▶ Many applicable dumb actions that do not progress towards our goal

Branching Factor

Forward Search

- ▶ **Forward search** can have a **very large** branching factor
 - ▶ Many applicable dumb actions that do not progress towards our goal
- ▶ The search algorithms we have covered can waste a lot of time here

Branching Factor

Forward Search

- ▶ **Forward search** can have a **very large** branching factor
 - ▶ Many applicable dumb actions that do not progress towards our goal
- ▶ The search algorithms we have covered can waste a lot of time here

Branching Factor


Forward Search

- ▶ **Forward search** can have a **very large** branching factor
 - ▶ Many applicable dumb actions that do not progress towards our goal
- ▶ The search algorithms we have covered can waste a lot of time here

It needs a good (domain-specific) heuristic or pruning procedure!


Backward search

How to plan?




1. Choose a **relevant action** that satisfies (some) goal propositions
2. Make a new **goal** by applying an **action** **backwards**:
 - ▶ **DEL** satisfied conditions of **goal**
 - ▶ **ADD** preconditions of **a**
 - ▶ Keep unsolved **goal** propositions
3. Repeat until the **goal** is satisfied by the start state

Backward search



What are the **relevant actions** here?

Backward search



What are the **relevant actions** here?

How do we know which one is the *best* one?

Branching factor

Backward search

- ▶ **Backward search** is not guided towards any **specific subgoal**

Branching factor

Backward search

- ▶ **Backward search** is not guided towards any **specific subgoal**
- ▶ The **order** in which we try to achieve the subgoals (and do search) matters

Branching factor

Backward search

- ▶ **Backward search** is not guided towards any **specific subgoal**
- ▶ The **order** in which we try to achieve the subgoals (and do search) matters
 - ▶ It impacts the efficiency of the search

Branching factor

Backward search

- ▶ **Backward search** is not guided towards any **specific subgoal**
- ▶ The **order** in which we try to achieve the subgoals (and do search) matters
 - ▶ It impacts the efficiency of the search
 - ▶ A wrong order can make the plan unfeasible

Sub-plans: Total and partial orders

How to plan

So far, we have only looked at algorithms that generate **complete** plans, with a **strict** order. However, there may exist **subplans**—sequence of actions that can be **partially** ordered.

Sub-plans: Total and partial orders

How to plan

So far, we have only looked at algorithms that generate **complete** plans, with a **strict** order. However, there may exist **subplans**—sequence of actions that can be **partially** ordered.

Think about it as when you put your socks and shoes on every morning:

Sub-plans: Total and partial orders

How to plan

So far, we have only looked at algorithms that generate **complete** plans, with a **strict** order. However, there may exist **subplans**—sequence of actions that can be **partially** ordered.

Think about it as when you put your socks and shoes on every morning:

- ▶ Which *sock* should go on first?

Sub-plans: Total and partial orders

How to plan

So far, we have only looked at algorithms that generate **complete** plans, with a **strict** order. However, there may exist **subplans**—sequence of actions that can be **partially** ordered.

Think about it as when you put your socks and shoes on every morning:

- ▶ Which *sock* should go on first?
- ▶ Do you do the subsequence $sock_1 \rightarrow shoe_1$ like a lunatic or $sock_1 \rightarrow sock_2$ first?

If we had enough hands I guess we could do $sock_i \rightarrow shoe_i$ in **parallel**!

Heuristic Planning

As usual, the **real world** is a much scarier place.

Heuristic Planning

As usual, the **real world** is a much scarier place.

- ▶ Problems (instances) can be huge (number of variables, number of constraints)

Heuristic Planning

As usual, the **real world** is a much scarier place.

- ▶ Problems (instances) can be huge (number of variables, number of constraints)
- ▶ We have limited resources / memory / time / money

Heuristic Planning

As usual, the **real world** is a much scarier place.

- ▶ Problems (instances) can be huge (number of variables, number of constraints)
- ▶ We have limited resources / memory / time / money
- ▶ We need faster solutions

Heuristic Planning

As usual, the **real world** is a much scarier place.

- ▶ Problems (instances) can be huge (number of variables, number of constraints)
- ▶ We have limited resources / memory / time / money
- ▶ We need faster solutions

Heuristic Planning

As usual, the **real world** is a much scarier place.

- ▶ Problems (instances) can be huge (number of variables, number of constraints)
- ▶ We have limited resources / memory / time / money
- ▶ We need faster solutions

We need **heuristic** solutions!

General idea: relax

Heuristic Planning

Most heuristics rely on **relaxing** the problem to make it easier.

General idea: relax

Heuristic Planning

Most heuristics rely on **relaxing** the problem to make it easier.

In planning:

- ▶ Add edges and **group** nodes (**subplans**)

General idea: relax

Heuristic Planning

Most heuristics rely on **relaxing** the problem to make it easier.

In planning:

- ▶ Add edges and **group** nodes (**subplans**)
 - ▶ State abstraction, pattern DBs, symmetry reduction...

General idea: relax

Heuristic Planning

Most heuristics rely on **relaxing** the problem to make it easier.

In planning:

- ▶ Add edges and **group** nodes (**subplans**)
 - ▶ State abstraction, pattern DBs, symmetry reduction...
- ▶ Ignore **restrictions**

General idea: relax

Heuristic Planning

Most heuristics rely on **relaxing** the problem to make it easier.

In planning:

- ▶ Add edges and **group** nodes (**subplans**)
 - ▶ State abstraction, pattern DBs, symmetry reduction. . .
- ▶ Ignore **restrictions**
 - ▶ Either **all** or **some** of them

General idea: relax

Heuristic Planning

Most heuristics rely on **relaxing** the problem to make it easier.

In planning:

- ▶ Add edges and **group** nodes (**subplans**)
 - ▶ State abstraction, pattern DBs, symmetry reduction. . .
- ▶ Ignore **restrictions**
 - ▶ Either **all** or **some** of them
- ▶ Ignore **negative fluents**

General idea: relax

Heuristic Planning

Most heuristics rely on **relaxing** the problem to make it easier.

In planning:

- ▶ Add edges and **group** nodes (**subplans**)
 - ▶ State abstraction, pattern DBs, symmetry reduction. . .
- ▶ Ignore **restrictions**
 - ▶ Either **all** or **some** of them
- ▶ Ignore **negative fluents**
- ▶ Weigh actions to have **preferred** actions

General idea: relax

Heuristic Planning

Most heuristics rely on **relaxing** the problem to make it easier.

In planning:

- ▶ Add edges and **group** nodes (**subplans**)
 - ▶ State abstraction, pattern DBs, symmetry reduction...
- ▶ Ignore **restrictions**
 - ▶ Either **all** or **some** of them
- ▶ Ignore **negative fluents**
- ▶ Weigh actions to have **preferred** actions
- ▶ Find **serialisable subplans**

General idea: relax

Heuristic Planning

Most heuristics rely on **relaxing** the problem to make it easier.

In planning:

- ▶ Add edges and **group** nodes (**subplans**)
 - ▶ State abstraction, pattern DBs, symmetry reduction. . .
- ▶ Ignore **restrictions**
 - ▶ Either **all** or **some** of them
- ▶ Ignore **negative fluents**
- ▶ Weigh actions to have **preferred** actions
- ▶ Find **serialisable subplans**
 - ▶ Achieving one will not undo the others.

General idea: relax

Heuristic Planning

Most heuristics rely on **relaxing** the problem to make it easier.

In planning:

- ▶ Add edges and **group** nodes (**subplans**)
 - ▶ State abstraction, pattern DBs, symmetry reduction. . .
- ▶ Ignore **restrictions**
 - ▶ Either **all** or **some** of them
- ▶ Ignore **negative fluents**
- ▶ Weigh actions to have **preferred** actions
- ▶ Find **serialisable subplans**
 - ▶ Achieving one will not undo the others.

General idea: relax

Heuristic Planning

Most heuristics rely on **relaxing** the problem to make it easier.

In planning:

- ▶ Add edges and **group** nodes (**subplans**)
 - ▶ State abstraction, pattern DBs, symmetry reduction...
- ▶ Ignore **restrictions**
 - ▶ Either **all** or **some** of them
- ▶ Ignore **negative fluents**
- ▶ Weigh actions to have **preferred** actions
- ▶ Find **serialisable subplans**
 - ▶ Achieving one will not undo the others.

...or try a **metaheuristic**!

Examples of relaxation

Heuristic Planning

- ▶ **State abstraction:** group subtasks into one bigger, more abstract task

Examples of relaxation

Heuristic Planning

- ▶ **State abstraction:** group subtasks into one bigger, more abstract task
 - ▶ All packages for Trondheim (from Oslo) count as one big package

Examples of relaxation

Heuristic Planning

- ▶ **State abstraction:** group subtasks into one bigger, more abstract task
 - ▶ All packages for Trondheim (from Oslo) count as one big package
 - ▶ Worry about getting those packages to Trondheim first. You can worry about individual deliveries afterwards!

Examples of relaxation

Heuristic Planning

- ▶ **State abstraction:** group subtasks into one bigger, more abstract task
 - ▶ All packages for Trondheim (from Oslo) count as one big package
 - ▶ Worry about getting those packages to Trondheim first. You can worry about individual deliveries afterwards!
 - ▶ Decompose costs: $Cost(P) = Cost(P_i) + Cost(P_j)$ where i, j are subgoals

Examples of relaxation

Heuristic Planning

- ▶ **State abstraction:** group subtasks into one bigger, more abstract task
 - ▶ All packages for Trondheim (from Oslo) count as one big package
 - ▶ Worry about getting those packages to Trondheim first. You can worry about individual deliveries afterwards!
 - ▶ Decompose costs: $Cost(P) = Cost(P_i) + Cost(P_j)$ where i, j are subgoals
- ▶ **Ignore restrictions:** 'a perfect plan *if* this road r were not closed'

Examples of relaxation

Heuristic Planning

- ▶ **State abstraction:** group subtasks into one bigger, more abstract task
 - ▶ All packages for Trondheim (from Oslo) count as one big package
 - ▶ Worry about getting those packages to Trondheim first. You can worry about individual deliveries afterwards!
 - ▶ Decompose costs: $Cost(P) = Cost(P_i) + Cost(P_j)$ where i, j are subgoals
- ▶ **Ignore restrictions:** 'a perfect plan *if* this road r were not closed'
 - ▶ We will try finding an alternative route for the r segment later!

Examples of relaxation

Heuristic Planning

- ▶ **State abstraction:** group subtasks into one bigger, more abstract task
 - ▶ All packages for Trondheim (from Oslo) count as one big package
 - ▶ Worry about getting those packages to Trondheim first. You can worry about individual deliveries afterwards!
 - ▶ Decompose costs: $Cost(P) = Cost(P_i) + Cost(P_j)$ where i, j are subgoals
- ▶ **Ignore restrictions:** ‘a perfect plan *if* this road r were not closed’
 - ▶ We will try finding an alternative route for the r segment later!
- ▶ **Serialisable subplans:** achieving a subgoal (putting on $shoe_1$) does not interfere with other goals (putting on $shoe_2$)

Classical Planning vs IRL

Planning in Complex Environments

So far, we have covered how to do **classical planning**:

Classical Planning vs IRL

Planning in Complex Environments

So far, we have covered how to do **classical planning**:

1. The environment is **fully observable**

Classical Planning vs IRL

Planning in Complex Environments

So far, we have covered how to do **classical planning**:

1. The environment is **fully observable**
 - ▶ We can “see” what happens at any state

Classical Planning vs IRL

Planning in Complex Environments

So far, we have covered how to do **classical planning**:

1. The environment is **fully observable**
 - ▶ We can “see” what happens at any state
2. The environment is **deterministic**

Classical Planning vs IRL

Planning in Complex Environments

So far, we have covered how to do **classical planning**:

1. The environment is **fully observable**
 - ▶ We can “see” what happens at any state
2. The environment is **deterministic**
 - ▶ Our actions always give us the same effect

Classical Planning vs IRL

Planning in Complex Environments

So far, we have covered how to do **classical planning**:

1. The environment is **fully observable**
 - ▶ We can “see” what happens at any state
2. The environment is **deterministic**
 - ▶ Our actions always give us the same effect
3. The environment is **static**

Classical Planning vs IRL

Planning in Complex Environments

So far, we have covered how to do **classical planning**:

1. The environment is **fully observable**
 - ▶ We can “see” what happens at any state
2. The environment is **deterministic**
 - ▶ Our actions always give us the same effect
3. The environment is **static**
 - ▶ It never changes, unless we act

Classical Planning vs IRL

Planning in Complex Environments

So far, we have covered how to do **classical planning**:

1. The environment is **fully observable**
 - ▶ We can “see” what happens at any state
2. The environment is **deterministic**
 - ▶ Our actions always give us the same effect
3. The environment is **static**
 - ▶ It never changes, unless we act

Classical Planning vs IRL

Planning in Complex Environments

So far, we have covered how to do **classical planning**:

1. The environment is **fully observable**
 - ▶ We can “see” what happens at any state
2. The environment is **deterministic**
 - ▶ Our actions always give us the same effect
3. The environment is **static**
 - ▶ It never changes, unless we act

How do we plan when we encounter more complex environments?

Planning in Complex Environments

1. I cannot see

Planning in Complex Environments

1. I cannot see

- ▶ So we make a **sensorless plan**

Planning in Complex Environments

1. I cannot see
 - ▶ So we make a **sensorless plan**
2. I can see but I am not sure what will happen

Planning in Complex Environments

1. I cannot see
 - ▶ So we make a **sensorless plan**
2. I can see but I am not sure what will happen
 - ▶ So we make a **contingency plan**

Planning in Complex Environments

1. I cannot see
 - ▶ So we make a **sensorless plan**
2. I can see but I am not sure what will happen
 - ▶ So we make a **contingency plan**
3. I can see and am sure what will happen but need to keep an eye

Planning in Complex Environments

1. I cannot see
 - ▶ So we make a **sensorless plan**
2. I can see but I am not sure what will happen
 - ▶ So we make a **contingency plan**
3. I can see and am sure what will happen but need to keep an eye
 - ▶ Then we need **online planning**

I cannot see

Sensorless planning

- ▶ I need to **make sure** that all **preconditions** are **met**
- ▶ I will then **carry out all operations** that will lead me to the **goal**

Example: paint a chair and a table with the same colour. How?

I Cannot See



I cannot see

Sensorless planning

- ▶ I need to **make sure** that all **preconditions** are **met**
- ▶ I will then **carry out all operations** that will lead me to the **goal**

Example: paint a chair and a table with the same colour. How?


[RemoveLid(Can), Paint(Chair, Can), Paint(Table, Can)]

I Cannot See




Recall

Belief space in sensorless planning



Recall

Belief space in non-deterministic world



- ▶ The agent knows where it is and see the dirt (if any) on its spot
- ▶ The **transition model** becomes a **function** of a **belief state**, an **action**, and a **another belief state**
 - ▶ In case of nondeterminism (right), we do like Dr. Strange and consider possible outcomes on different universes. **How?**

I can see but I am not sure what will happen

Contingency planning

- ▶ I need to **make sure** to know **where I am** by **looking around**
- ▶ Then, and depending on **where I am**, I will **carry out** necessary operations **conditionally**

Example: paint a chair and a table with the same colour. How?



I can see but I am not sure what will happen

Contingency planning

- ▶ I need to **make sure** to know *where I am* by **looking around**
- ▶ Then, and depending on *where I am*, I will **carry out** necessary operations **conditionally**

Example: paint a chair and a table with the same colour. How?



FullPlan =


*[LookAt(Table), LookAt(Chair),
if Color(Table, c) \wedge Color(Chair, c), then NoOP
else ContingencyPlan]*

I can see but I am not sure what will happen

Contingency planning

- ▶ I need to **make sure** to know **where I am** by **looking around**
- ▶ Then, and depending on **where I am**, I will **carry out** necessary operations **conditionally**

Example: paint a chair and a table with the same colour. How?



ContingencyPlan =


```
[RemoveLid(Can), LookAt(Can),  
if Color(Table, c)  $\wedge$  Color(Can, c) then Paint(Chair, can)  
else if Color(Chair, c)  $\wedge$  Color(can, c) then Paint(Table, can)  
else[Paint(Chair, Can), Paint(Table, Can)]]
```

I need to keep an eye

Online planning

The world is **dynamic**, and it can change in unpredictable ways.

- ▶ **Action monitoring**: check that all **preconditions** hold
- ▶ **Plan monitoring**: check if the **plan** can **succeed** (or the **goal** is **satisfiable**)
- ▶ **Goal monitoring**: check if there is a *better* **goal**



The Job-shop Scheduling Problem



Figure 1




Figure 2

Machine	O1	O2	O3	O4
J1	1	4	2	3
J2	2	3	1	4
J3	3	4	2	1
J4	4	1	2	3

Table 1 Machine Sequence

Time	O1	O2	O3	O4
J1	3	2	4	3
J2	4	5	2	2
J3	4	4	3	1
J4	3	4	1	1

Table 2 Processing Time

Image from Ataç, 2023: [Job Shop Scheduling Problem and Solution Algorithms](#)

JSP: heuristic approach

The Job-shop Scheduling Problem

- ▶ Assign the **earliest** and **latest** possible start for each **action**

JSP: heuristic approach

The Job-shop Scheduling Problem

- ▶ Assign the **earliest** and **latest** possible start for each **action**
- ▶ Then duration of the plan is that of the **critical path**

JSP: heuristic approach

The Job-shop Scheduling Problem

- ▶ Assign the **earliest** and **latest** possible start for each **action**
- ▶ Then duration of the plan is that of the **critical path**
 - ▶ i.e., the path with the **longest duration**. There cannot be a shorter plan than this!

JSP: heuristic approach

The Job-shop Scheduling Problem

- ▶ Assign the **earliest** and **latest** possible start for each **action**
- ▶ Then duration of the plan is that of the **critical path**
 - ▶ i.e., the path with the **longest duration**. There cannot be a shorter plan than this!
- ▶ The JSP is an \mathcal{NP} -problem

JSP: heuristic approach

The Job-shop Scheduling Problem

- ▶ Assign the **earliest** and **latest** possible start for each **action**
- ▶ Then duration of the plan is that of the **critical path**
 - ▶ i.e., the path with the **longest duration**. There cannot be a shorter plan than this!
- ▶ The JSP is an \mathcal{NP} -problem
 - ▶ The optimisation version (finding how) is an \mathcal{NP} -hard problem

JSP: heuristic approach

The Job-shop Scheduling Problem

- ▶ Assign the **earliest** and **latest** possible start for each **action**
- ▶ Then duration of the plan is that of the **critical path**
 - ▶ i.e., the path with the **longest duration**. There cannot be a shorter plan than this!
- ▶ The JSP is an \mathcal{NP} -problem
 - ▶ The optimisation version (finding how) is an \mathcal{NP} -hard problem

JSP: heuristic approach

The Job-shop Scheduling Problem

- ▶ Assign the **earliest** and **latest** possible start for each **action**
- ▶ Then duration of the plan is that of the **critical path**
 - ▶ i.e., the path with the **longest duration**. There cannot be a shorter plan than this!
- ▶ The JSP is an \mathcal{NP} -problem
 - ▶ The optimisation version (finding how) is an \mathcal{NP} -hard problem

In practice we use **metaheuristics**!