

TDT4171 Artificial Intelligence Methods

Lecture 7 – Learning from Observations

Norwegian University of Science and Technology

Helge Langseth
Gamle Fysikk 255
helge.langseth@ntnu.no



- 1 Summary from last time
- 2 Learning from observations
 - Learning agents
 - Inductive learning
 - Measuring learning performance
 - Overfitting
- 3 Introduction to ANNs
 - Background
 - Perceptrons
 - Gradient descent
- 4 Deep Learning
 - Representation
 - Regularization

Summary from last time



- **Sequential decision problems**
 - **Assumptions:** Stationarity, Markov assumption, Additive rewards, infinite horizon with discount
 - **Model classes:** Markov decision processes/Partially Observable Markov Decision Processes
 - **Algorithm:** Value iteration / policy iteration
- Intuitively, MDPs combine **probabilistic models over time** (filtering, prediction) with **maximum expected utility principle**.

Reference Group meeting:

We'll have a meeting in the RefGrp next week. Send them an email if you have feedback.

Learning



This is the second part of the course:

- We have learned about **representations** for uncertain knowledge
- **Inference** in these representations
- Making **decisions** based on the inferences

Now we will talk about **learning** the representations:

- Supervised learning – When focus is on learning “a mapping”
- Reinforcement learning – When focus is on learning “to behave”
- There is a third category of learning, unsupervised learning, that we won't cover in this course.

Plan for Part II of the course



Today: Recap: Machine learning, neural nets, deep nets

March 7th: **Instance based learning and CBR.**

Guest lecture, Kerstin Bach

March 14th: More deep learning

March 21st: (Deep) Reinforcement Learning

March 28th: NLP including RLHF, and Transformers

April 4th: Summary

April 11th: Class trip. **No lecture**

April 18th: Easter. **No lecture**

April 25th: Buffer. Hopefully **No lecture**

Assignments: We keep releasing them according to current schedule. Last assignment (Assignment 10) has planned deadline April 3rd. 7 out of 10 needed.

Learning goals for today



Being familiar with:

- **Motivation** and **Formalization** for learning
- **Fundamental issues** like learning bias, overfitting
- **Neural net** basics and main idea for learning
- **Deep Learning** basics

Recap: Learning from observations

Why do Learning?



Arthur Samuel playing Checkers, 1956

Well-defined learning problem



What “parts” are needed to formally define a learning problem?
... and how can Arthur Samuel describe his problem in that way?

Discuss with your neighbour for a couple of minutes.

Well-defined learning problem



What “parts” are needed to formally define a learning problem?
... and how can Arthur Samuel describe his problem in that way?

One classically relates a learning-problem to three objects: **Task** T , **Performance measure** P , and **Experience** E :

- Improve over task T :
 - Playing checkers.
- ... with respect to performance measure P :
 - Games (out of 100, say) won against a fixed opponent.
- ... based on experience E :
 - Playing against itself to generate experience to learn from.

Why do Learning?



- Learning **modifies the agent's decision mechanisms** to improve performance
- Learning is essential for **unknown environments** (when designer lacks omniscience)
- Learning is useful as a **system construction** method (expose the agent to reality rather than trying to write it down)

Currently we see lots of work in machine learning, due to:

- Availability of data massively increasing
- Increased utilization of hardware architectures (like GPUs)
- Method development (like deep learning)
- Clever definition of new tasks as machine learning problems

Inductive learning



Simplest form: Learn a function from examples

f is the **target function**

An **example** is a pair $\{x, f(x)\}$, e.g., $\left\{ \begin{array}{c|c|c} O & O & X \\ \hline & X & \\ \hline X & & \end{array} , +1 \right\}$

Problem:

Find **hypothesis** $h \in H$ s.t. $h \approx f$ given a **training set** of examples

This is a highly simplified model of real learning:


- Ignores **prior knowledge**
- Assumes a **deterministic, observable** “environment”
- Assumes examples are **given**

Inductive learning method



- Construct/adjust h to agree with f on training set
- h is **consistent** if it agrees with f on all examples

Example – curve fitting:




Inductive learning method



- Construct/adjust h to agree with f on training set
- h is **consistent** if it agrees with f on all examples

Example – curve fitting:




Inductive learning method



- Construct/adjust h to agree with f on training set
- h is **consistent** if it agrees with f on all examples

Example – curve fitting:




Inductive learning method



- Construct/adjust h to agree with f on training set
- h is **consistent** if it agrees with f on all examples

Example – curve fitting:




Which curve is better? – and WHY?
Can we make an operational definition?
Discuss with your neighbour for a couple of minutes.

Inductive learning method



- Construct/adjust h to agree with f on training set
- h is **consistent** if it agrees with f on all examples

Example – curve fitting:



Ockham's razor: maximize consistency and simplicity.

Key insight: We don't necessarily aim to be consistent.

Rather, we typically aim to make good predictions!


Performance measurement



Question: How do we know that $h \approx f$?

Answer: Try h on a new **test set** of examples (use **same distribution over example space** as training set)

Learning curve = % correct on **test set** as a function of training set size




Performance measurement contd.



Learning curve depends on...

- **realizable** (can express target function) vs. **non-realizable**
- **non-realizability** can be due to missing attributes or restricted hypothesis class (e.g., thresholded linear function)
- **redundant expressiveness** (e.g., loads of irrelevant attributes)



Overfitting



Consider error of hypothesis h over

- Training data: $\text{error}_t(h)$
- Entire distribution \mathcal{D} of data (often approximated by measurement on test-set): $\text{error}_{\mathcal{D}}(h)$

Overfitting

Hypothesis $h \in H$ **overfits** training data if there is an alternative hypothesis $h' \in H$ such that

$$\text{error}_t(h) < \text{error}_t(h') \text{ and } \text{error}_{\mathcal{D}}(h) > \text{error}_{\mathcal{D}}(h')$$

Avoiding Overfitting



- Overfitting often occur for flexible learning representations (like high-order polynomials, ...)
- Overfitting harms the usefulness of the machine learning system, because it is the **generalization ability** (score on a test-set) that is important!

What techniques can be used to prevent overfitting for ML in general (not specific to particular learning algorithm)?

Discuss with your neighbour for a couple of minutes.

Avoiding Overfitting



- Overfitting often occur for flexible learning representations (like high-order polynomials, ...)
- Overfitting harms the usefulness of the machine learning system, because it is the **generalization ability** (score on a test-set) that is important!

What techniques can be used to prevent overfitting for ML in general (not specific to particular learning algorithm)?

- Compare models' actual **generalization ability** using a validation set or some statistical technique on training data.
- Use some heuristics, like bias model search towards **simplicity** (e.g., linear over high order polynomial fit)

Recap: Basics of neural nets

Connectionist Models



Some facts about the human brain:

- Number of neurons about 10^{11}
- Connections per neuron about $10^4 - 10^5$
- Scene recognition time about .1 second
- Neuron switching time about .001 second
- 100 inference steps doesn't seem like enough for scene recognition → much parallel computation

Properties of artificial neural nets (ANN's):


- Many neuron-like threshold switching units
- Many weighted interconnections among units
- Highly parallel, distributed process
- Emphasis on tuning weights automatically

Brains



10^{11} neurons of > 20 types, 10^{14} synapses


Signals are noisy “spike trains” of electrical potential



Perceptron




Output is a **nonlinear function** of the inputs with offset:




A gross **oversimplification** of real neurons, but its purpose is to develop understanding of what networks of simple units can do.

Activation functions



(a)



(b)

(a) is a **step function** or **threshold function**

(b) is a **sigmoid** function $1/(1 + e^{-x})$

(c) We will also consider **linear** functions and **rectified linear units** (ReLU).

Note! Changing the bias b shifts the output along the x -axis

Finding the optimal weights



Let's start with a **linear unit**, where input is \vec{x} gives output

$$o = \underbrace{w_0 \cdot \overbrace{x_0}^{x_0=1}}_{\text{per convention: } b} + w_1 x_1 + \cdots + w_n x_n = \vec{w}^\top \vec{x}$$

We learn w_i 's that minimise some **loss**. For regression models it makes sense to use the squared error

$$\mathcal{L}[\vec{w}] = \sum_{d \in \mathcal{D}} (t_d - o_d)^2,$$

where \mathcal{D} is set of **training examples**, each of the form $d = \langle \mathbf{x}_d, t_d \rangle$.

Finding optimal weights – requirements



Description of our situation:

- We have a function $\mathcal{L}[\vec{w}]$ we want to minimise (wrt \vec{w}).
- **Why not just try a number of weight configurations \vec{w}_i , calculate $\mathcal{L}[\vec{w}_i]$ and see what happens?**
- There are infinitely many (even uncountably many) weight configurations.
- Minimization is typically in very high dimensional space.
- Evaluating $\mathcal{L}[\vec{w}]$ involves summing over all training examples – can be very expensive.

We cannot use a standard trial & error approach, but must devise a local search method. **What can we do instead?**

Discuss with your neighbour for a couple of minutes.

Gradient Descent – The setup



- We want to find the value x which minimizes $f(x)$.
Yes, $f(x)$ will be replaced by $\mathcal{L}[\vec{w}]$ later on!
- To avoid evaluating the whole function we use an iterative approach:
 - Guess a value for x
 - Calculate the derivative at x .
 - Make a new guess for x based on the calculated information
 - ... and keep going.
- **Intuition:**
 - If the derivative is zero then we are done
 - If it is small (in absolute value) we are close to the minimizing point
 - If it is large we are not that close


Solution:

Use update rule $x_{i+1} \leftarrow x_i - \eta \cdot f'(x_i)$. $\eta > 0$ is the **learning rate**.

Gradient Descent – Example



Minimize $f(x) = 2x^4 - 5x^3 + 2x^2 - 6x + 4$ with $\eta = 0.025$.




The $f(x)$ has a minimum at $x = 1.8261$. Let's try to find it...

Gradient Descent – Example



Minimize $f(x) = 2x^4 - 5x^3 + 2x^2 - 6x + 4$ with $\eta = 0.025$.




Starting from $x_0 = 3$ and finding $f'(3) = 87$:

$$\begin{aligned}x_1 &= x_0 - \eta f'(x_0) \\ &= 3 - 0.025 \cdot 87 = 0.8250\end{aligned}$$

Gradient Descent – Example



Minimize $f(x) = 2x^4 - 5x^3 + 2x^2 - 6x + 4$ with $\eta = 0.025$.




Going from $x_1 = 0.8250$ with $f'(0.8250) = -8.4172$:

$$\begin{aligned}x_2 &= x_1 - \eta f'(x_1) \\ &= 0.825 - 0.025 \cdot (-8.4172) = 1.0354\end{aligned}$$

Gradient Descent – Example



Minimize $f(x) = 2x^4 - 5x^3 + 2x^2 - 6x + 4$ with $\eta = 0.025$.




Going from $x_2 = 1.0354$ with $f'(1.0354) = -9.0592$:

$$x_3 = 1.0354 - 0.025 \cdot (-9.0592) = 1.2619$$

Gradient Descent – Example



Minimize $f(x) = 2x^4 - 5x^3 + 2x^2 - 6x + 4$ with $\eta = 0.025$.




... and finally going from $x_{10} = 1.8260$ with $f'(1.8260) = -0.0034$:

$$x_{11} = 1.8260 - 0.025 \cdot (-0.0034) = 1.8261$$

... and we are done.

Gradient Descent – in higher dimensions



Recall that


- The **gradient** of a surface $\mathcal{L}[\vec{w}]$ is a vector in the direction the curve grows the most (calculated at \vec{w}).
- The gradient is calculated as $\nabla \mathcal{L}[\vec{w}] \equiv \left[\frac{\partial \mathcal{L}}{\partial w_0}, \frac{\partial \mathcal{L}}{\partial w_1}, \dots, \frac{\partial \mathcal{L}}{\partial w_n} \right]$.

$$\text{Training rule: } \Delta \vec{w} = -\eta \cdot \nabla \mathcal{L}[\vec{w}], \text{ i.e., } \Delta w_i = -\eta \cdot \frac{\partial \mathcal{L}}{\partial w_i}.$$

The perceptron's problem: Expressibility




Output from perceptron on input (x_1, x_2)



- Adjusting weights moves the location, orientation, and steepness of cliff
- Cannot tackle “correlation effects” of non-separable targets
- Solution: Make **layers** of nodes. **All continuous functions representable w/ 1 hidden layer, all functions w/ 2**

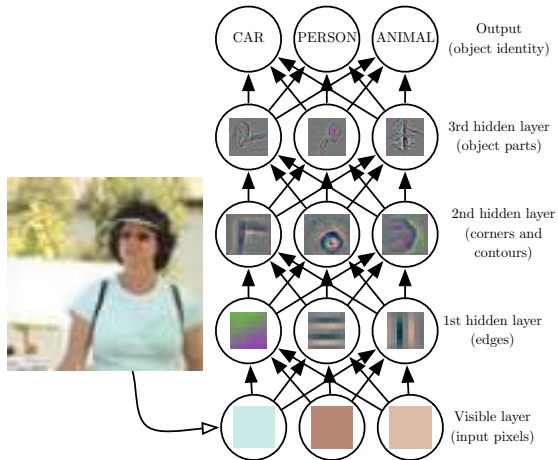
Recap: Deep Learning basics

Representation matters




- Representation in cartesian coordinates (x and y) is “difficult”: Not linearly separable.
- Representation in polar coordinates (r and θ) is “easy”.
But how do we find this representation?

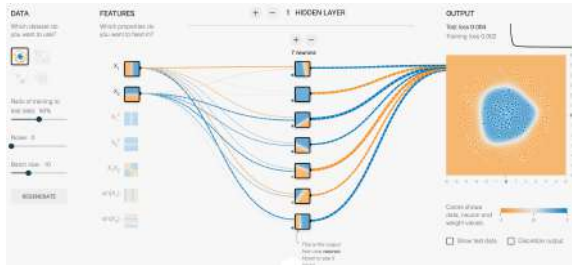
Depth: Compositions repeated



Depth: Compositions repeated



Tensorflow playground



<https://playground.tensorflow.org/>

Convolutional Neural nets



- **Goal:** Scale up to process very large images/videos
 - Sparse connections
 - Parameter sharing
 - Automatically generalize across spatial translations of inputs
 - Applicable to any input laid out on a grid (1-D, 2-D, 3-D, ...) and other data with spatial structure
- **Key idea:** Replace/replicate flattened representations and matrix multiplication with convolution that respect locality of information!
 - Everything else stays the same
 - Optimization criteria
 - Training algorithm
 - And so on

Convolutions




Mathematical definition: $(f_1 * f_2)(t) = \int_{-\infty}^{\infty} f_1(\tau) \cdot f_2(t - \tau) d\tau$

Convolutions




Mathematical definition: $(f_1 * f_2)(t) = \int_{-\infty}^{\infty} f_1(\tau) \cdot f_2(t - \tau) d\tau$



Convolutions




Mathematical definition: $(f_1 * f_2)(t) = \int_{-\infty}^{\infty} f_1(\tau) \cdot f_2(t - \tau) d\tau$




Convolutional Neural nets: Efficiency




Limited “spread”:



Reuse of params:



Implementation of DL models



Implementation frameworks:

- **tensorflow.keras:** High level of abstraction, easy to get going, used in enterprise systems and cloud platforms.
- **pytorch:** More coding required, far easier to customize for semi-advanced extensions.

Model types:

- **Feed forward:** Simple but ineffective
- **Conv.nets:** Parameter efficient; similar (slightly better) quality

Expressive Capabilities of ANNs



Boolean functions:

- Every boolean function can be represented by network with a single hidden layer
- Note: We might require exponential (in number of inputs) hidden units

Continuous functions:

- Every bounded continuous function can be approximated with arbitrarily small error, by network with one hidden layer
- Any function can be approximated to arbitrary accuracy by a network with two hidden layers

Depth adds to expressiveness

- Overfitting can be a massive issue
- Techniques for **regularization** becomes very important

Regularization in deep learning



Remember the def of overfitting

Hypothesis $h \in H$ **overfits** training data if there is an alternative hypothesis $h' \in H$ such that

$$\text{error}_t(h) < \text{error}_t(h') \text{ and } \text{error}_{\mathcal{D}}(h) > \text{error}_{\mathcal{D}}(h')$$

And now: Regularization

Regularization is any modification we make to a learning algorithm that is intended to reduce its generalization error but not its training error.

Types of regularization:

- Norm penalty
- Dropout
- ...

Norm penalties






- Remember our objective: Optimize some loss $\mathcal{L}_{\text{data}}(\mathbf{w})$ where \mathbf{w} are the weights in the deep net
- Norm penalty (a.k.a. *weight regularization*) penalizes “long \mathbf{w} -vectors”:

$$\mathcal{L}_{\text{norm}}(\mathbf{w}) = \|\mathbf{w}\|_p = \left(\sum_j w_j^p \right)^{1/p}$$

- Total loss: $\mathcal{L}(\mathbf{w}) = \mathcal{L}_{\text{data}}(\mathbf{w}) + \lambda \cdot \mathcal{L}_{\text{norm}}(\mathbf{w})$. Note the λ , which balances the two losses.
- Typical examples:
 - $p = 1$: Encourages sparsity (some weights “exactly” zero)
 - $p = 2$: Typically results in small but non-zero weights

Norm penalties




$$\mathcal{L}(w) = \mathcal{L}_{\text{data}}(w) + \lambda \cdot \|w\|_p$$

Dropout



Dropout: To randomly drop a fraction of neurons (or sometimes also edges) at each training iteration.



Summary



- **Learning:** Use data to generate (or improve) a representation.
- **Inductive learning hypothesis:** “Old data has relevance for new problems”
- **Neural networks:** Capable structures defined by combining many simple computational units
 - Simple perceptrons
 - Layered models → Feed forward networks
 - DL models can also contain more complicated structures
- **Learning:** Define a loss, minimize using gradient descent
- **Overfitting:** A model overfits if it does well on the training data, but fails to generalize
 - Prefer simplicity
 - Can be enforced using, e.g., norm penalties; dropout

Next week: Guest lecture by Kerstin Bach: CBR