

TDT4171 Artificial Intelligence Methods

Lecture 9 – More on Deep Learning

Norwegian University of Science and Technology

Helge Langseth
Gamle Fysikk 255
helge.langseth@ntnu.no



- 1 Summary from last time
- 2 Modelling with neural nets
- 3 Gradient-based learning
 - General setup
 - Perceptrons
 - Backpropagation
- 4 Fancy representations
 - Convolutions
 - RNNs
- 5 Summary



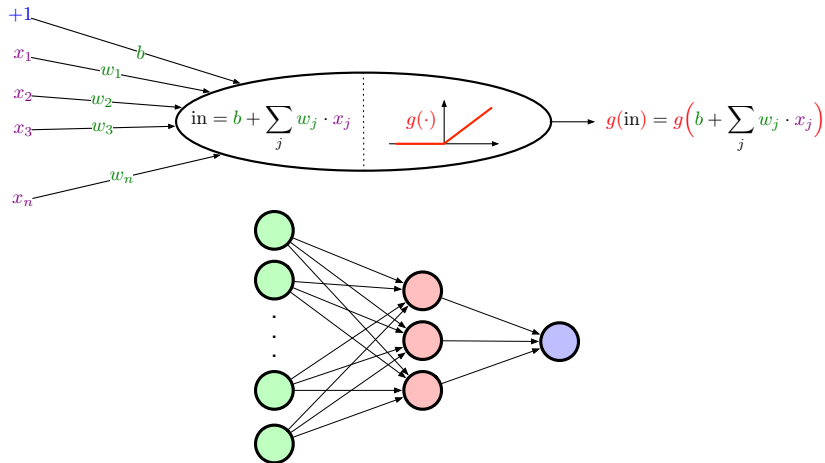
Summary from last time

- **Learning:** Use data to generate (or improve) a representation.
- **Inductive learning hypothesis:** “Old data has relevance for new problems”
- **Neural networks:** Capable structures defined by combining many simple computational units
 - Simple perceptrons
 - Layered models → Feed forward networks
 - DL models can also contain more complicated structures
- **Learning:** Define a loss, minimize using gradient descent
- **Overfitting:** A model overfits if it does well on the training data, but fails to generalize
 - Prefer simplicity
 - Can be enforced using, e.g., norm penalties; dropout

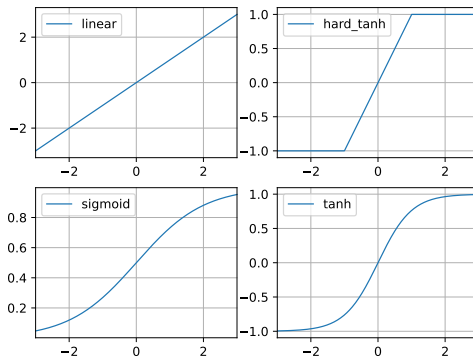
Reference group meeting:

- Piazza
- Assignment deadlines
- Time-keeping

Model structure: Layers of simple nodes

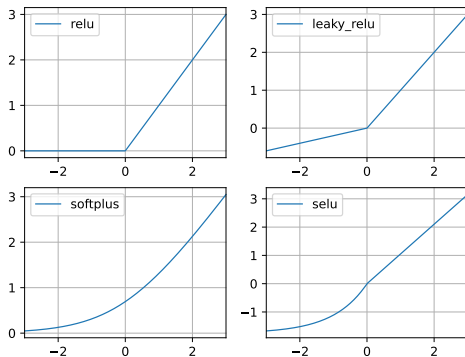


Model structure: Activations for hidden nodes



- Classic activation functions, chosen for simplicity (“linear”, “hard_tanh”) or biological plausibility (“sigmoid”) – the latter has a sibling that also generates negative values (“tanh”).
- Sigmoids no longer in much use (for hidden nodes) due to diminishing gradients.

Model structure: Activations for hidden nodes



- The “relu” (with extensions) very popular nowadays.
- Gradients zero \rightarrow “leaky_relu”; Derivative undefined at zero \rightarrow “softplus”; Size explosion \rightarrow “selu”.
- Finding new activation functions is still an active research field.

Model structure: Activations for output nodes



- **Regression problems:** Depending on output range, typically “linear” (for unbounded), “relu” (for positive) or “sigmoid”/“tanh” with proper scaling (for ranged)
- **Classification problems:** Nice if output layer is a probability distribution (non-negative values, sum to 1).
 - Assume we have K classes, and $\mathbf{in} = (\mathbf{in}_1, \dots, \mathbf{in}_K)$ be pre-activation values. Then

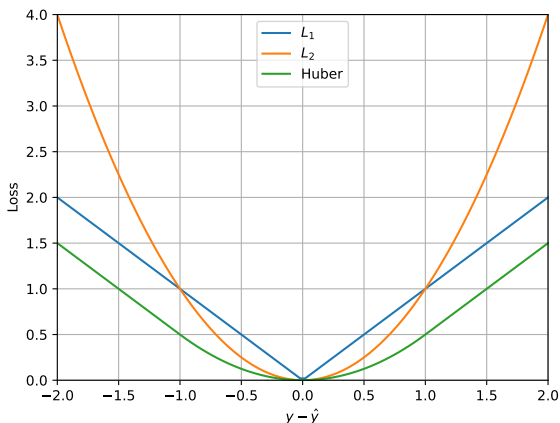
$$\text{softmax}(\mathbf{in})_i = \frac{\exp(\mathbf{in}_i)}{\sum_{j=1}^K \exp(\mathbf{in}_j)}.$$

- No constraints on $(\mathbf{in}_1, \dots, \mathbf{in}_K)$ wrt. range of values:
 $\mathbf{in} \in \mathbb{R}^K \Rightarrow \text{softmax}(\mathbf{in})_i \geq 0, \sum_{k=1}^K \text{softmax}(\mathbf{in})_k = 1.$
- Defines a “competition” between output nodes (i.e., classes).

Model structure: Loss functions



- **Regression problems:** Typically L_p -loss: $\mathcal{L}(y, \hat{y}) = |y - \hat{y}|^p$ for $p = 1$ or 2 (or Huber-loss; a combo).



Model structure: Loss functions

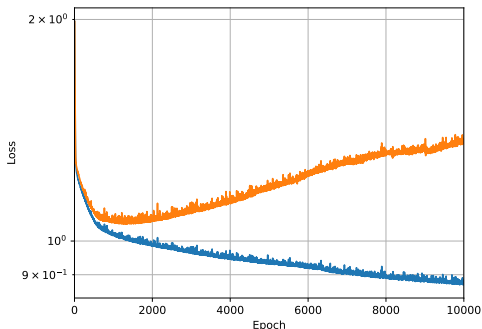


- **Regression problems:** Typically L_p -loss: $\mathcal{L}(y, \hat{y}) = |y - \hat{y}|^p$ for $p = 1$ or 2 (or Huber-loss; a combo).
- **Classification problems:** Typically cross entropy loss $\mathcal{L}(\mathbf{y}, \hat{\mathbf{y}}) = -\sum_{k=1}^K y_k \cdot \log_2(\hat{y}_k)$. Here, \mathbf{y} is a one-hot-encoding over the classes (a vector of length K , with all zeros except the position for the correct class, which is one).

Predictions	$\hat{\mathbf{y}}$:	0.20	0.25	0.55
One-hot encoding	\mathbf{y} :	0	1	0
	$-y_k \log_2(\hat{y}_k)$:	0	2	0
Total loss (in bits): 2.0				

Predictions	$\hat{\mathbf{y}}$:	.02	.05	.93
One-hot encoding	\mathbf{y} :	0	0	1
	$-y_k \log_2(\hat{y}_k)$:	0	0	0.1
Total loss (in bits): 0.1				

Model structure: Regularization



- Typical example of **overfitting**: Training loss (blue) keeps going down, yet validation loss (orange) starts increasing.
- Obvious strategy: **Early stopping**.
- Additional problem: Gap between training and validation loss indicate need for **regularization**.

Model structure: Regularization (cont'd)



Definition of Regularization:

Reduction of **testing** error while maintaining a low **training** error.

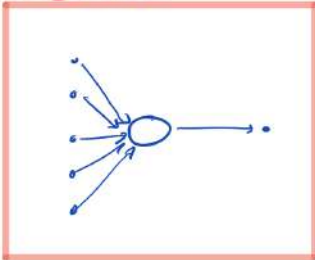
Motivation:

- Excessive training does reduce training error, but often at the expense of higher testing error.
- Flexible model classes (like neural nets) are prone to this.
- By effectively **memorizing** training-examples, we do not (necessarily) promote the ability to **generalize**

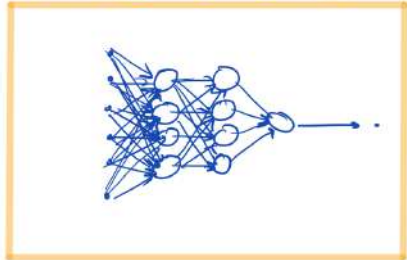
Types of regularization :

- Early stopping
- Weight-norm regularization (see slide from last week)
- Dropout (see slide from last week)
- ... and lots more!

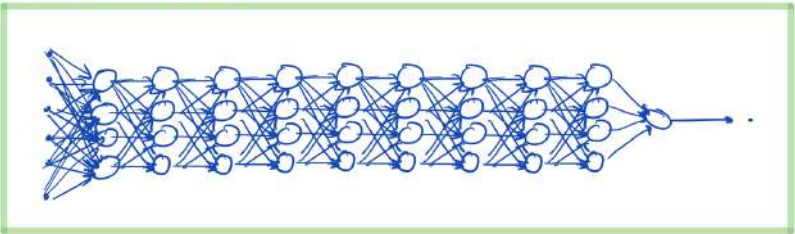
PERCEPTRON



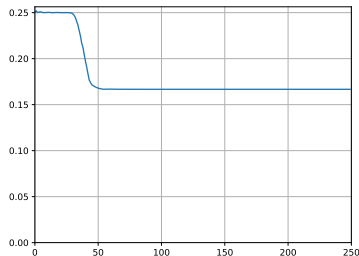
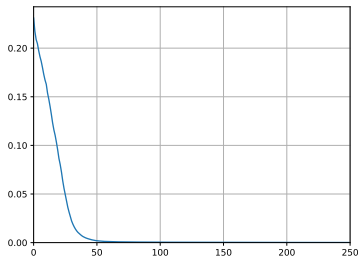
"STANDARD NEURAL NET"



"DEEP LEARNING"



Does it work?



XOR on a small network

- One epoch: All possible XOR configs.
- Model is $2 \rightarrow 2 \rightarrow 1$. Small, so hard to learn.
- All weights randomly initialized in $[-1, 1]$.
- Sometimes learning fails, in this case initialization was crucial.

How to make it work

- ① You need to be willing to spend a lot of time on hyper-parameter tuning. Tools are available. Use them!
- ② Be restrictive on model size.
Too big models take longer to learn and may overfit more.
- ③ Don't be too restrictive on model size.
Too small models may be unable to represent relationship.
- ④ Scale all input values to the same range.
- ⑤ Be careful when initializing weights: Start small!
- ⑥ Regularize
- ⑦ Monitor *training loss* and *validation loss*.
The former should become very small, while the latter should be similar and not increase.

Gradient-based learning

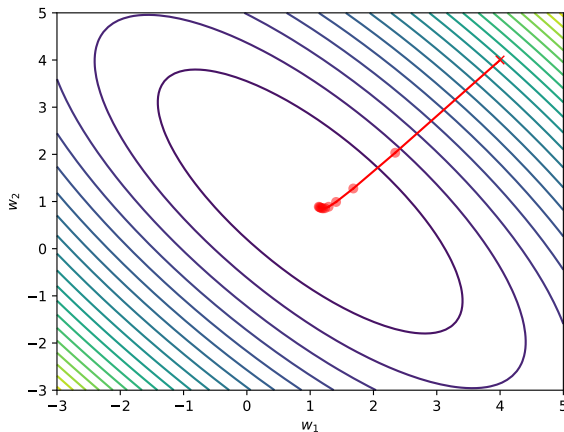


Minimize the the objective function \mathcal{L} using **partial derivatives** of \mathcal{L} wrt. model weights \mathbf{w} !

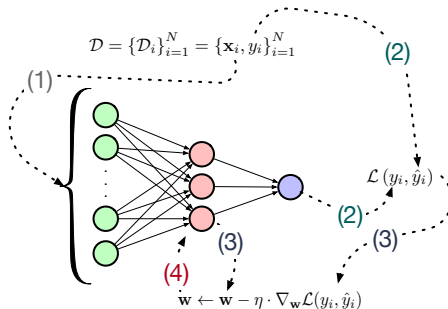
Simple setup: Learn the function $y = f(\mathbf{x})$:

- Data: $\mathcal{D}_i = (\mathbf{x}_i, y_i)$; Model weights: \mathbf{w} ; Output: $\hat{y} = \mathbf{w}^\top \mathbf{x}$.
- Objective: $\mathcal{L}(y, \hat{y}) = (y - \hat{y})^2 = (y - \mathbf{w}^\top \mathbf{x})^2$.
- **Note!** We will not see $y = f(\mathbf{x})$ for **all** possible \mathbf{x} , but will optimize based on our data $\{\mathcal{D}_i\}_{i=1}^N$. Loss is a sum over all datapoints.
- Since $\frac{\partial}{\partial w_j} \mathcal{L}(y, \mathbf{w}^\top \mathbf{x}) = -2x_j \cdot (y - \mathbf{w}^\top \mathbf{x})$, the move based on (\mathbf{x}, y) would be
$$w_j \leftarrow w_j - \eta \cdot \frac{\partial}{\partial w_j} \mathcal{L}(y, \mathbf{w}^\top \mathbf{x}) = w_j - \eta \cdot 2x_j \cdot (\mathbf{w}^\top \mathbf{x} - y)$$
- Full dataset: $w_j \leftarrow w_j - \eta \cdot \sum_{i=1}^N \frac{\partial}{\partial w_j} \mathcal{L}(y_i, \mathbf{w}^\top \mathbf{x}_i)$

Gradient-based learning



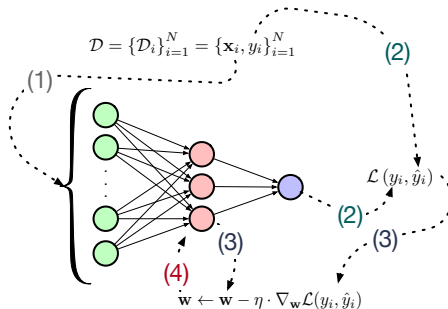
Gradient-based learning in neural nets



Steps:

- ➊ Push \mathbf{x}_i **forwards** to calculate \hat{y}_i .
- ➋ Calculate **loss** based on \hat{y}_i and observed y_i : $\mathcal{L}(y_i, \hat{y}_i)$.
- ➌ Calculate gradients $\nabla_{\mathbf{w}} \mathcal{L}(y_i, \hat{y}_i)$ while moving **backwards**.
- ➍ Update **weights**.

Gradient-based learning in neural nets



This is general purpose!

- NNs are getting increasingly complex, but the general idea is the same: Forward pass to find \hat{y} , backward pass to calculate $\nabla_{\mathbf{w}} \mathcal{L}(y_i, \hat{y}_i)$.
- Finding the gradients can be tedious and challenging, but frameworks like Tensorflow and PyTorch help us.

Gradient Descent – Details for our simple example



Recall our simple example:

Simple setup: Learn the function $f(x)$:

- Data: $\mathcal{D}_i = (\mathbf{x}_i, y_i)$; Model weights: \mathbf{w} ; Output: $\hat{y} = \mathbf{w}^\top \mathbf{x}$.
- Objective: $\mathcal{L}(y, \hat{y}) = (y - \hat{y})^2 = (y - \mathbf{w}^\top \mathbf{x})^2$.
- This corresponds to a **perceptron** w/ **linear** activation.
- We have already seen the result:

$$w_j \leftarrow w_j - \eta \cdot \frac{\partial}{\partial w_j} \mathcal{L}(y, \mathbf{w}^\top \mathbf{x}) = w_j - \eta \cdot 2x_j \cdot (\mathbf{w}^\top \mathbf{x} - y)$$

- This is the general rule for a perceptron with identity transfer.

Gradient Descent – Details for our simple example



$$\begin{aligned}\frac{\partial \mathcal{L}}{\partial w_j} &= \frac{\partial}{\partial w_j} \sum_{i=1}^N (y_i - \hat{y}_i)^2 \\&= \sum_{i=1}^N \frac{\partial}{\partial w_j} (y_i - \hat{y}_i)^2 \\&= \sum_{i=1}^N 2(y_i - \hat{y}_i) \frac{\partial}{\partial w_j} (y_i - \hat{y}_i) \\&= 2 \sum_{i=1}^N (y_i - \hat{y}_i) \frac{\partial}{\partial w_j} (y_i - \mathbf{w}^\top \mathbf{x}_i) \\&= -2 \sum_{i=1}^N (y_i - \hat{y}_i) \left(\frac{\partial}{\partial w_j} \sum_l w_l \cdot x_{i,l} \right) \\ \frac{\partial \mathcal{L}}{\partial w_j} &= -2 \sum_{i=1}^N x_{i,j} \cdot (y_i - \hat{y}_i)\end{aligned}$$

Gradient Descent – The algorithm for simple example



Gradient-Descent(\mathcal{D}, η)

Each training example is a pair of the form $\langle \mathbf{x}, y \rangle$, where \mathbf{x} is the vector of input values, and y is the target output value. η is the learning rate (e.g., .05).

- ① Initialize each w_j to some small random value
- ② Until the termination condition is met:
 - ① Initialize: $\Delta w_j \leftarrow 0$.
 - ② For each $\langle \mathbf{x}_i, y_i \rangle$ in \mathcal{D} :
 - Send \mathbf{x}_i through the network and compute the output \hat{y}_i
 - For each linear unit weight w_j : $\Delta w_j \leftarrow \Delta w_j - 2(y_i - \hat{y}_i) \cdot x_{i,j}$
 - ③ For each linear unit weight w_j : $w_j \leftarrow w_j - \eta \cdot \Delta w_j$

Same algorithm skeleton works for other activation functions, as long as we can calculate $\frac{\partial \mathcal{L}}{\partial w_j}$.

Gradient Descent: Perceptron with transfer functions



Assume the logistic transfer function:

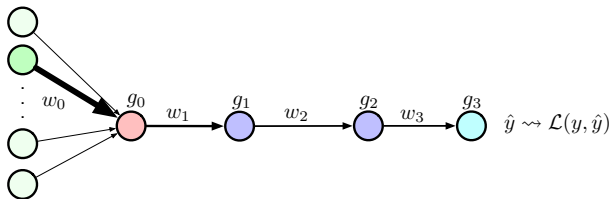
$$g(t) = \frac{1}{1 + \exp(-t)}, \quad g'(t) = g(t) \cdot [1 - g(t)].$$

$$\begin{aligned} \frac{\partial \mathcal{L}}{\partial w_j} &= \frac{\partial}{\partial w_j} \sum_{i=1}^N (y_i - \hat{y}_i)^2 \\ &= 2 \sum_{i=1}^N (y_i - \hat{y}_i) \frac{\partial}{\partial w_j} (y_i - g(\mathbf{w}^\top \mathbf{x}_i)) \\ &= 2 \sum_{i=1}^N (y_i - \hat{y}_i) \cdot (-1) \cdot x_{i,j} \cdot g'(\mathbf{w}^\top \mathbf{x}_i) \end{aligned}$$

$$\frac{\partial \mathcal{L}}{\partial w_j} = -2 \sum_{i=1}^N x_{i,j} \cdot (y_i - \hat{y}_i) \cdot \underbrace{\hat{y}_i \cdot (1 - \hat{y}_i)}$$

This is new from the “simple example”

Backpropagation Algorithm – Single chain

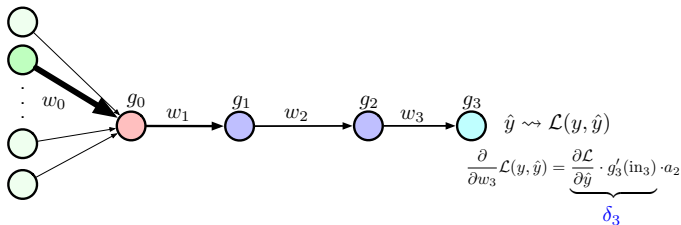


Notation for node j :

- Weight into node is w_j ; its activation function $g_j(\cdot)$.
- in_j as the weighted input.
- a_j is output from node j , $a_j = g_j(\text{in}_j)$ and $\text{in}_{j+1} = w_{j+1} \cdot a_j = w_{j+1} \cdot g_j(\text{in}_j)$.

Goal: Calculate $\nabla_{\mathbf{w}} \mathcal{L}(y, \hat{y}) = \left[\frac{\partial \mathcal{L}}{\partial w_0}, \frac{\partial \mathcal{L}}{\partial w_1}, \frac{\partial \mathcal{L}}{\partial w_2}, \frac{\partial \mathcal{L}}{\partial w_3} \right]^\top$.

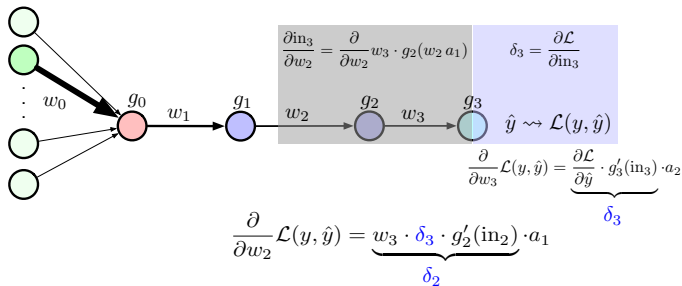
Backpropagation Algorithm – Single chain



Calculation:

- Start from back: $\frac{\partial \mathcal{L}}{\partial w_3} = \frac{\partial \mathcal{L}}{\partial \hat{y}} \cdot g'_3(\text{in}_3) \cdot a_2$; $\text{in}_3 = w_3 a_2$.
- $\delta_3 \leftarrow \frac{\partial \mathcal{L}}{\partial \text{in}_3} = \frac{\partial \mathcal{L}}{\partial \hat{y}} \cdot g'_3(\text{in}_3)$. $\frac{\partial \mathcal{L}}{\partial w_3} = \frac{\partial \mathcal{L}}{\partial \text{in}_3} \frac{\partial \text{in}_3}{\partial w_3} = \delta_3 a_2$.

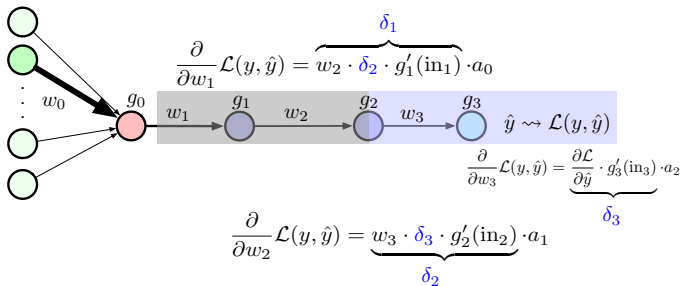
Backpropagation Algorithm – Single chain



Calculation:

- Start from back: $\frac{\partial \mathcal{L}}{\partial w_3} = \frac{\partial \mathcal{L}}{\partial \hat{y}} \cdot g'_3(\text{in}_3) \cdot a_2$; $\text{in}_3 = w_3 a_2$.
 $\delta_3 \leftarrow \frac{\partial \mathcal{L}}{\partial \text{in}_3} = \frac{\partial \mathcal{L}}{\partial \hat{y}} \cdot g'_3(\text{in}_3)$. $\frac{\partial \mathcal{L}}{\partial w_3} = \frac{\partial \mathcal{L}}{\partial \text{in}_3} \frac{\partial \text{in}_3}{\partial w_3} = \delta_3 a_2$.
- $\frac{\partial \mathcal{L}}{\partial w_2} = \frac{\partial \mathcal{L}}{\partial \text{in}_3} \frac{\partial \text{in}_3}{\partial w_2} = \delta_3 \cdot w_3 \frac{\partial}{\partial w_2} g_2(w_2 \cdot a_1) = \delta_3 \cdot w_3 \cdot g'_2(\text{in}_2) \cdot a_1$;
 $\delta_2 \leftarrow \frac{\partial \mathcal{L}}{\partial \text{in}_2} = \delta_3 \cdot w_3 \cdot g'_2(\text{in}_2)$.

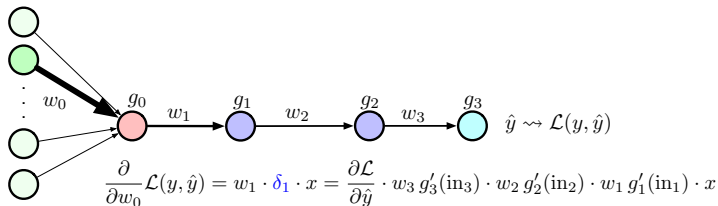
Backpropagation Algorithm – Single chain



Calculation:

- Start from back: $\frac{\partial \mathcal{L}}{\partial w_3} = \frac{\partial \mathcal{L}}{\partial \hat{y}} \cdot g'_3(\text{in}_3) \cdot a_2$; $\text{in}_3 = w_3 a_2$.
 $\delta_3 \leftarrow \frac{\partial \mathcal{L}}{\partial \text{in}_3} = \frac{\partial \mathcal{L}}{\partial \hat{y}} \cdot g'_3(\text{in}_3)$. $\frac{\partial \mathcal{L}}{\partial w_3} = \frac{\partial \mathcal{L}}{\partial \text{in}_3} \frac{\partial \text{in}_3}{\partial w_3} = \delta_3 a_2$.
- $\frac{\partial \mathcal{L}}{\partial w_2} = \frac{\partial \mathcal{L}}{\partial \text{in}_3} \frac{\partial \text{in}_3}{\partial w_2} = \delta_3 \cdot w_3 \frac{\partial}{\partial w_2} \cdot g_2(w_2 \cdot a_1) = \delta_3 \cdot w_3 \cdot g'_2(\text{in}_2) \cdot a_1$;
 $\delta_2 \leftarrow \frac{\partial \mathcal{L}}{\partial \text{in}_2} = \delta_3 \cdot w_3 \cdot g'_2(\text{in}_2)$.
- $\frac{\partial \mathcal{L}}{\partial w_1} = \frac{\partial \mathcal{L}}{\partial \text{in}_2} \frac{\partial \text{in}_2}{\partial w_1} = \delta_2 \cdot w_2 \cdot g'_1(\text{in}_1) \cdot a_0$.

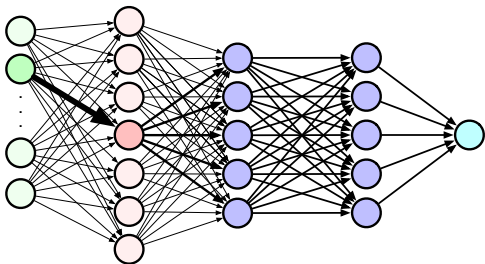
Backpropagation Algorithm – Single chain



Notice:

- **Forward** for activations: $a_{j+1} = g_{j+1}(w_{j+1} a_j)$.
- **Backward** for derivatives:
 - $\delta_j = \frac{\partial \mathcal{L}}{\partial \text{in}_j} = g'_j(\text{in}_j) w_{j+1} \cdot \delta_{j+1}$.
 - $\frac{\partial \mathcal{L}}{\partial w_j} = \frac{\partial \mathcal{L}}{\partial \text{in}_j} \frac{\partial \text{in}_j}{\partial w_j} = \delta_j \cdot a_{j-1}$

Backpropagation Algorithm – General model



Layered model:

- The effect of changing w_0 will spread through many additive “paths”; all must be considered when adapting the weight.
- **The derivatives of a sum equals the sum of derivatives**, so this is actually “easy” to generalize.
- Now $\delta_j = \frac{\partial \mathcal{L}}{\partial \text{in}_j}$ is a **vector** with one element per node in layer j :

$$\delta_{j,k} = g'_j(\text{in}_{j,k}) \sum_{\ell} w_{j+1,\ell} \cdot \delta_{j+1,\ell}.$$

Backpropagation Algorithm – Sigmoid units

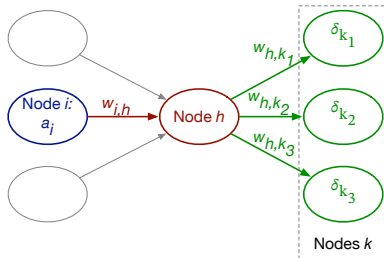


- ❶ Initialize all weights to small random numbers.
- ❷ Until satisfied:
 - For each training example
 - ❶ Input the training example to the network and compute the network outputs and node activations along the way
 - ❷ For each output unit k : $\delta_k \leftarrow \underbrace{-2(y_k - \hat{y}_k)}_{\text{Loss contrib.}} \times \underbrace{\hat{y}_k(1 - \hat{y}_k)}_{g'(\text{in}_k)}$
 - ❸ For each hidden unit h :

$$\delta_h \leftarrow \underbrace{a_h(1 - a_h)}_{g'(\text{in}_h)} \times \underbrace{\sum_{k \in \text{outputs}} w_{h,k} \cdot \delta_k}_{\text{Unit } h\text{'s contribution to next layer's error}}$$

- ❹ Update each network weight $w_{i,j}$: $w_{i,j} \leftarrow w_{i,j} - \eta \Delta w_{i,j}$, where $\Delta w_{i,j} = \delta_j \cdot a_i$ and a_i is the activation of the node pushing information into $w_{i,j}$.

Backprop in yet another picture

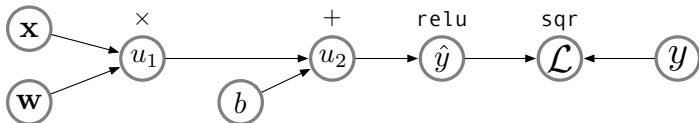


$$\delta_h \leftarrow a_h(1 - a_h) \times \sum_k w_{h,k} \delta_k \quad \text{and} \quad \Delta w_{i,h} \leftarrow \delta_h \cdot a_i$$

Note the “locality” of the calculations:

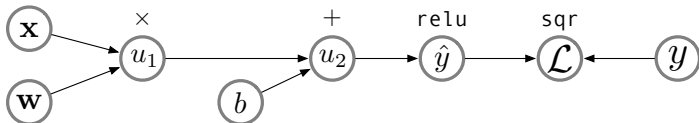
- $\Delta w_{i,h}$ only cares about Node h , its parent (Node i), and its children (Nodes k).
- The calculation of $\delta_h = \frac{\partial \mathcal{L}}{\partial \text{in}_h}$ does not care about the parent (Node i), so we can use the same value for all parents!

Backprop in DL: The computational graph



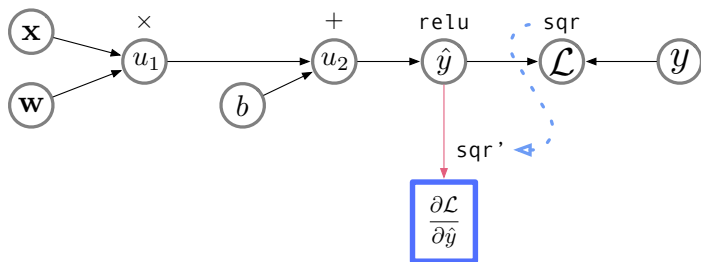
- Representation to describe computations in a directed graph. This one shows $\hat{y} = \text{relu}(\mathbf{x}^\top \mathbf{w} + b)$; $\mathcal{L} = (y - \hat{y})^2$.
- The layered NN architecture maps directly to the comp.graph:
 - All calculation results are nodes
 - All terms used in a node's calculations are parents.
- Gradients are found using the graph and the **chain rule**.

Backprop in DL: The computational graph

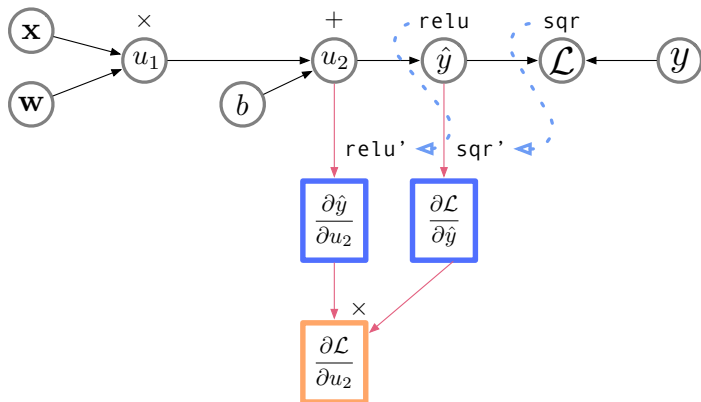


- Calculating derivatives by going backwards in the graph:
 - Do a local compute at node
 - Multiply with result from child.
 - If several children: Just sum
- Starting at \mathcal{L} and ending at b gives us $\frac{\partial \mathcal{L}}{\partial b}$; starting at \mathcal{L} and ending at w produces $\frac{\partial \mathcal{L}}{\partial w}$.

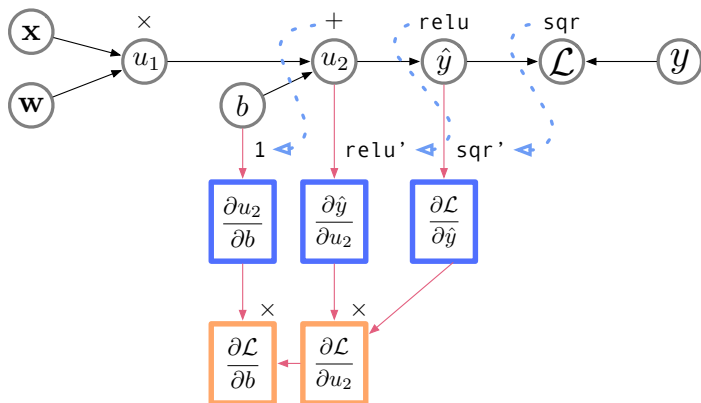
Backprop in DL: The computational graph



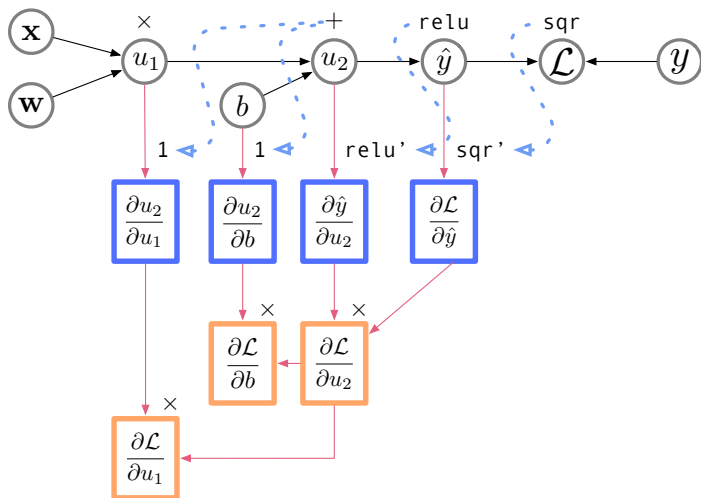
Backprop in DL: The computational graph



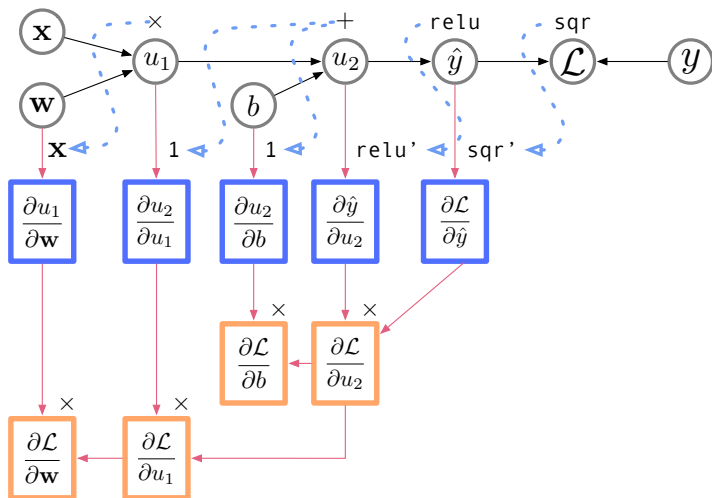
Backprop in DL: The computational graph



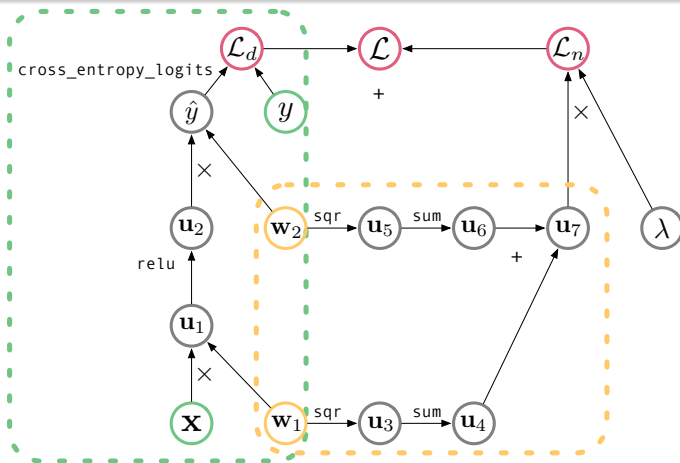
Backprop in DL: The computational graph



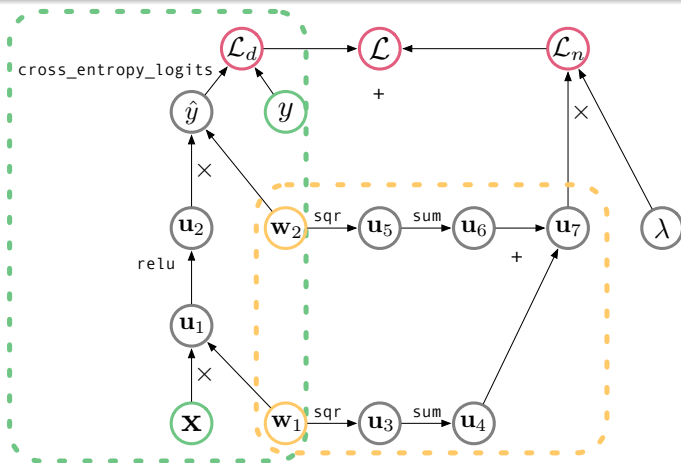
Backprop in DL: The computational graph



Computational graph for classification



Computational graph for classification



$$\mathcal{L} = \text{cross_entropy_logits}(y, w_2^T \underbrace{\text{relu}(w_1^T x))}_{u_1}) + \underbrace{\lambda \cdot (\|w_1\|_2 + \|w_2\|_2)}_{u_7}$$

Convergence of Backpropagation



Gradient descent to some local minimum:

- Perhaps not the *global* minimum ...
 - Include weight *momentum* $\alpha > 0$:

$$\Delta w_{i,j}(n) = \left(\underbrace{2\delta_j \cdot a_i}_{\text{The "standard"}} + \underbrace{\alpha \cdot \Delta w_{i,j}(n-1)}_{\text{Scaled last move}} \right)$$

- Train multiple nets with different initial weights
- Batching/stochastic gradient descent

Nature of convergence – depending on g (here: logistic):

- Initialize weights near zero
→ Therefore, initial networks *near-linear*.
- Increasingly non-linear functions possible as training progresses.

Backpropagation summary



- Gradient descent over entire *network* weight vector
- Easily generalized to arbitrary directed graphs
- Will find a local, not necessarily global error minimum
 - In practice, often works well
- Minimizes error over *training* examples
 - Will it generalize well to subsequent examples? **Overfitting...**
- Training can take thousands of iterations → **slow!**
- Using network after training is very fast

Convolutional Neural nets

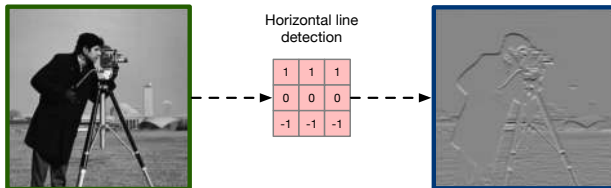
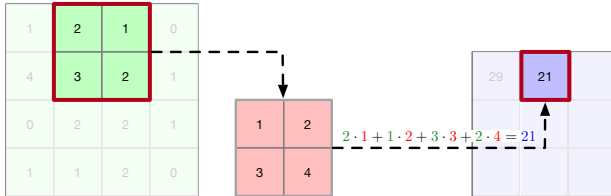


- **Goal:** Scale up to process very large images/videos
 - Sparse connections
 - Parameter sharing
 - Automatically generalize across spatial translations of inputs
 - Applicable to any input laid out on a grid (1-D, 2-D, 3-D, ...) and other data with spatial structure
- **Key idea:** Replace/replicate flattened representations and matrix multiplication with convolution that respect locality of information. **Everything else stays the same:**
 - Optimization criteria
 - Training algorithm
 - And so on ...

Convolutions



Std. definition: $(f_1 * f_2)(t) = \int_{-\infty}^{\infty} f_1(\tau) \cdot f_2(t - \tau) d\tau$



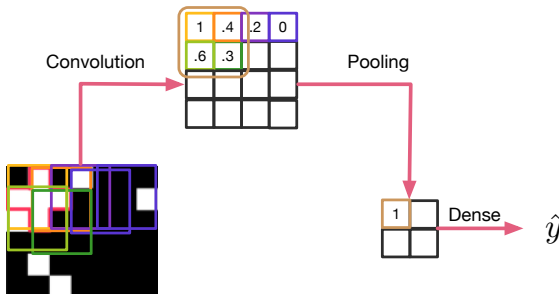
Learning convolutions



Data: Random binary images, some with a cross



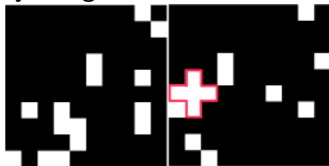
Model:



Learning convolutions



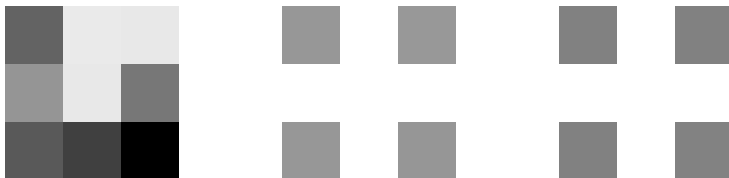
Data: Random binary images, some with a cross



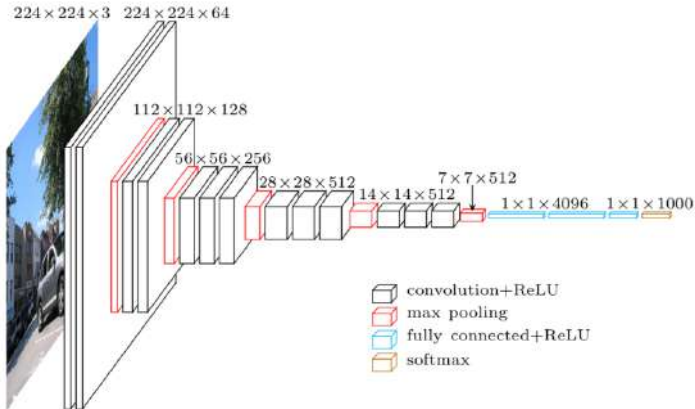
Model:

Single 3×3 convolutional kernel, one filter, with `relu`,
max-pooling down to 2×2 , then flatten and logistic on output.

Results: Kernel after 1, 10, and 20 epochs:



Classic CNN model architecture: VGG16 from 2014

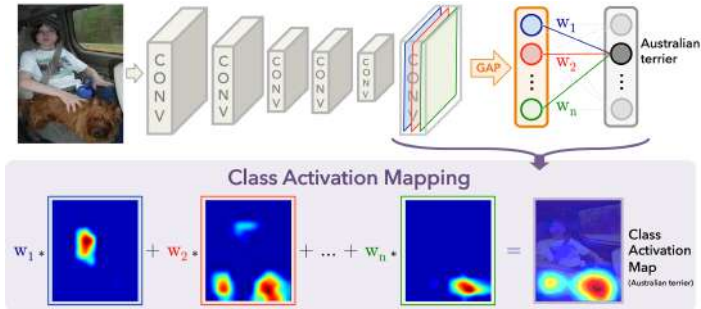


XAI: Understanding what the model does



Early example: Class Activation Mapping (CAM, 2015)

- After last convolution: Average over spatial positions, so each filter is represented by one variable
- Learn classifier over that \rightarrow weights per filter and class.
- Use location-based firing combined with filter weights to find class activation map.



Recurrent Neural Networks (RNNs)



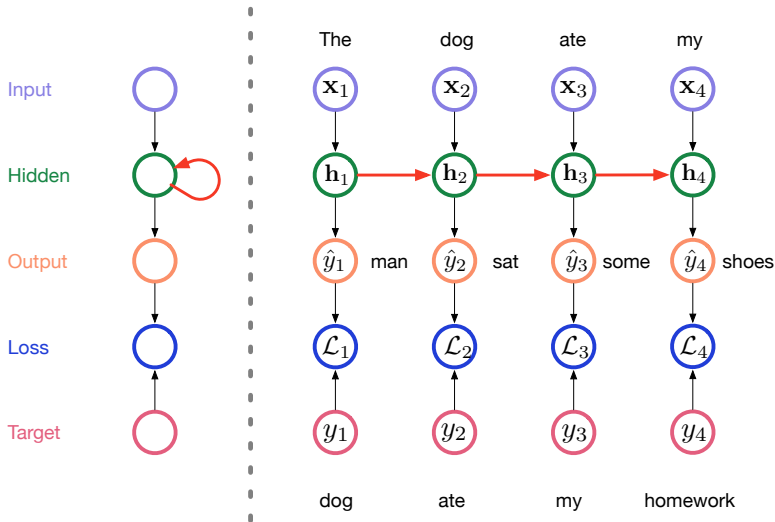
Typical use-cases for RNNs

- Inputs can be variable-length sequences; order is important; long-term dependencies.; parameter sharing/stationarity.
- Analysis of time-series data (e.g., measurements stock tickers over time) and predictions over those
- Sequential data, in particular anything with language: Modeling, Generation, Translation, Recognition, ...
- **Transformers** (attention-based models) and **conv-nets** are getting increasingly popular in this domain.

RNNs

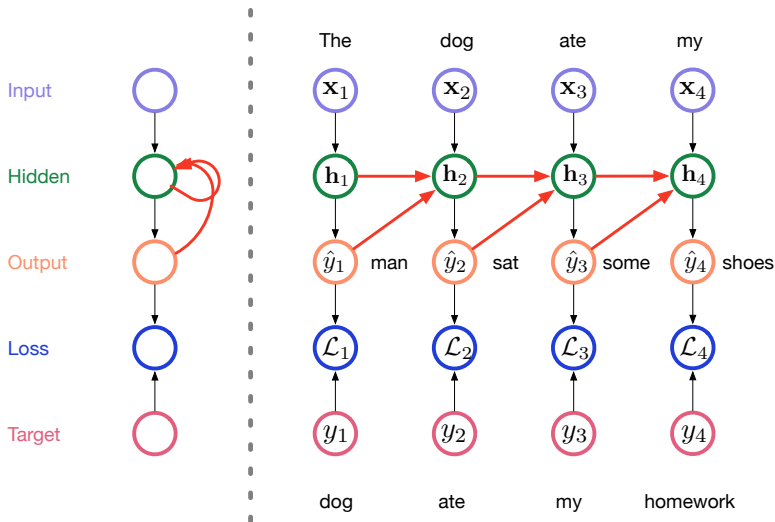
- A neural network that models sequential data (as HMMs)
- Many versions exist; sometimes tailor-made to accommodate particular signal flows (resembling DBN modelling task).

Standard architectures: Here for next-word prediction



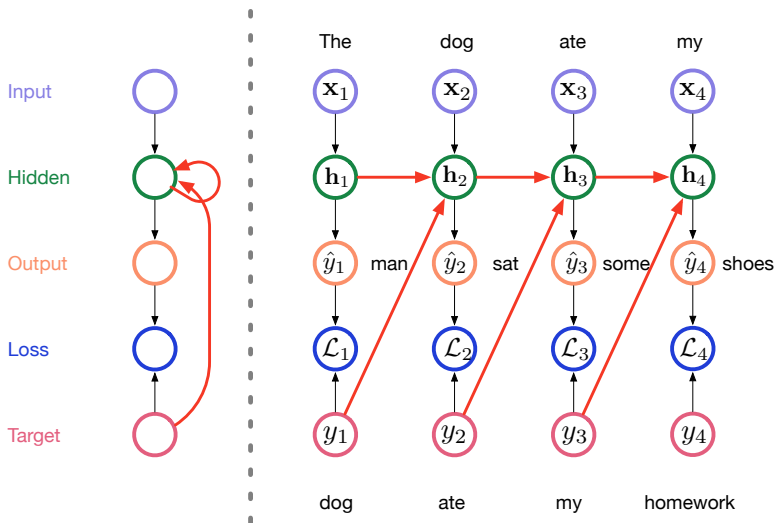
Plan: HMM-like structure

Standard architectures: Here for next-word prediction



Plan: Push generated tokens

Standard architectures: Here for next-word prediction



Plan: Push *correct* tokens during training; generated during test.

Standard architectures: Here for next-word prediction



- Other architectures (beyond “one-to-one-repeated”) exist, like
 - **Many-to-one:** Sequence classification
 - **One-to-many:** Captioning
 - **Sequence-to-sequence:** Translation
- **Problematic issue:** Track long-term information?

My dog disappeared. I went by my mom's house searching, had a coffee, and met my sister. She has moved away, and I don't see her that much. Then I went home, talked to my wife, and watched a ballgame. My team won. I had a good beer, one of my favourites, that is now available at the local grocery store. Then I found him! **Question:** Who was found?

- Size of h_t may have to extremely large to transport all that is needed to know. Hard to learn!
- Can we somehow **enforce structure** on the model to simplify information aggregation?
 - Yes, using LSTMs! (Next slide)
 - Yes, using transformers! (in two weeks)

Long Short Term Memory/LSTM



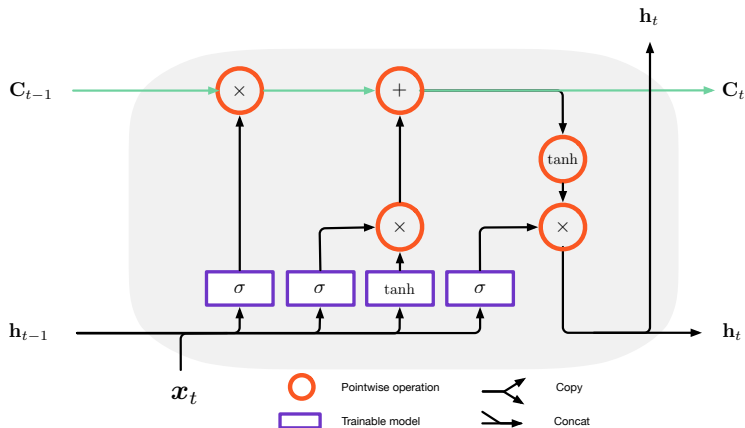
LSTM background:

- Invented by Hochreiter and Schmidhuber in 1997.
- Lots of different versions; ; massively popular since 2009.
- Can be combined with other layers (e.g., convolution)

Modelling components:

- Cell-state C_t and hidden state h_t . Both passed through time.
- $C_t \leftarrow \alpha_C([h_{t-1}, x_t]) \cdot C_{t-1} + \beta_C([h_{t-1}, x_t])$: Modulate based on input and hidden, add based on input and hidden;
- $h_t \leftarrow \alpha_h([h_{t-1}, x_t]) \cdot \tanh(C_t)$: Modulate transformed C_t using input and hidden; h_t is also the output at time t .

LSTMs (Hochreiter&Schmidhuber, 1997)



LSTM gradients need not be diminishing
 → Can learn long-term relationships.

Summary



- **Deep learning** is the most prominent ML technique nowadays.
- **Main reasons for success:** Data, compute, algorithmic developments.
- Model tuning using **gradient-based** techniques. Backprop is the general-purpose workhorse.
- **Modelling** is still important: CNNs, RNNs incl. LSTMs, (and later: Transformers) are defined to incorporate specific biases.